

Report Submission

Student ID: 23701011

Student Full Name: Nilanjana Das Jui

Academic Semester: 5th Semester

Session: 2022-2023

Batch: 23

Email: nilanjana.csecu@gmail.com

Department name: Computer Science and Engineering

Faculty Name: Engineering

Hall: Bijoy 24

University name: University of Chittagong

Course Code: CSE - 512

Course Title: Operating System Lab

Course Credits: 2

Date of Submission: February 05, 2026

Contents

1 Experiment 1: UNIX Permission and umask Calculator	8
1.1 Problem Statement	8
1.2 Solution Implementation	8
1.3 Sample Input 1	11
1.4 Sample Output 1	12
1.5 Sample Input 2	12
1.6 Sample Output 2	12
1.7 Sample Input 3	12
1.8 Sample Output 3	12
2 Experiment 2: POSIX File Copy with open/read/write	13
2.1 Problem Statement	13
2.2 Solution Implementation	13
2.3 Sample Input 1	18
2.4 Sample Output 1	18
2.5 Sample Input 2	18
2.6 Sample Output 2	19
3 Experiment 3: Directory Listing and Metadata Report (ls + stat subset)	20
3.1 Problem Statement	20
3.2 Solution Implementation	20

3.3	Sample Input 1: Directory does not exist	24
3.4	Sample Output 1	24
3.5	Commands to create test directory and files	25
3.6	Sample Input 2: List directory sorted by size after creation	25
3.7	Sample Output 2	25
3.8	Sample Input 3: Directory anotherdir	25
3.9	Sample Output 3	26
3.10	Sample Input 4: Directory missingdir (does not exist)	26
3.11	Sample Output 4	26
4	Experiment 4: grep-lite: Deterministic Text Pattern Search	27
4.1	Problem Statement	27
4.2	Solution Implementation	27
4.3	Commands to create test files	31
4.4	Sample Input 1: Files not yet created	31
4.5	Sample Output 1: Files missing	31
4.6	Sample Input 2: After creating test files	31
4.7	Sample Output 2: No matches	32
4.8	Sample Input 3: Pattern "hello"	32
4.9	Sample Output 3: One match	32
5	Experiment 5: Process Spawner and Exit-Status Reporter (fork/exec/wait)	33
5.1	Problem Statement	33

5.2	Solution Implementation	33
5.3	Sample Input 1	36
5.4	Sample Output 1	36
5.5	Sample Input 2: Invalid command	36
5.6	Sample Output 2	37
6	Experiment 6: Signal-Based Timeout Supervisor (sigaction + alarm + kill)	38
6.1	Problem Statement	38
6.2	Solution Implementation	38
6.3	Sample Input 1: Command completes before timeout	41
6.4	Sample Output 1	42
6.5	Sample Input 2: Command exceeds timeout	42
6.6	Sample Output 2	42
6.7	Sample Input 3	42
6.8	Sample Output 3	42
7	Experiment 7: Pipe-Based Filter Chain (pipe + dup2)	43
7.1	Problem Statement	43
7.2	Solution Implementation	43
7.3	Sample Input 1	46
7.4	Sample Output 1	46
7.5	Sample Input 2: Producer fails	47
7.6	Sample Output 2	47

8 Experiment 8: Shared Memory Counter IPC (shm_open + mmap + sem_open)	48
8.1 Problem Statement	48
8.2 Solution Implementation	48
8.3 Sample Input 1	51
8.4 Sample Output 1	52
8.5 Sample Input 2: Invalid Arguments	52
8.6 Sample Output 2	52
9 Experiment 9: Threaded Deterministic Reducer (pthreads + mutex)	53
9.1 Problem Statement	53
9.2 Solution Implementation	53
9.3 Sample Input 1	55
9.4 Sample Output 1	55
9.5 Sample Input 2: Invalid Arguments	55
9.6 Sample Output 2	55
10 Experiment 10: Bounded Buffer Producer-Consumer with Semaphores	56
10.1 Problem Statement	56
10.2 Solution Implementation	56
10.3 Sample Input 1	59
10.4 Sample Output 1	59
10.5 Sample Input 2: Invalid Arguments	59
10.6 Sample Output 2	59

11 Experiment 11: CPU Scheduling Simulator I (FCFS and Non-preemptive SJF)	60
11.1 Problem Statement	60
11.2 Solution Implementation	60
11.3 Sample Input	63
11.4 Sample Output	64
 12 Experiment 12: CPU Scheduling Simulator II (Round Robin)	65
12.1 Problem Statement	65
12.2 Solution Implementation	65
12.3 Sample Input	68
12.4 Sample Output	69
 13 Experiment 13: Priority Scheduling Simulator (Non-preemptive with Aging)	70
13.1 Problem Statement	70
13.2 Solution Implementation	70
13.3 Sample Input	73
13.4 Sample Output	74
 14 Experiment 14: Deadlock Avoidance using Banker's Algorithm	75
14.1 Problem Statement	75
14.2 Solution Implementation	75
14.3 Sample Input	78
14.4 Sample Output	79

15 Experiment 15: Deadlock Detection via Wait-For Graph Cycle	80
15.1 Problem Statement	80
15.2 Solution Implementation	80
15.3 Sample Input 1	84
15.4 Sample Output 1	84
15.5 Sample Input 2	84
15.6 Sample Output 2	84
 16 Experiment 16: Contiguous Memory Allocation Simulator (First/Best/-Worst Fit)	 85
16.1 Problem Statement	85
16.2 Solution Implementation	85
16.3 Sample Input	89
16.4 Sample Output	89
 17 Experiment 17: Paging Address Translation with Optional TLB	 90
17.1 Problem Statement	90
17.2 Solution Implementation	91
17.3 Sample Input	94
17.4 Sample Output	94
 18 Experiment 18: Page Replacement Simulator (FIFO, LRU, OPT)	 95
18.1 Problem Statement	95
18.2 Solution Implementation	95

18.3 Sample Input	100
18.4 Sample Output	100
19 Experiment 19: File Allocation Strategy Simulator (Contiguous, Linked, Indexed)	101
19.1 Problem Statement	102
19.2 Solution Implementation	102
19.3 Sample Input	108
19.4 Sample Output	108
20 Experiment 20: Disk Scheduling Simulator (FCFS, SSTF, SCAN, C-SCAN)	109
20.1 Problem Statement	109
20.2 Solution Implementation	110
20.3 Sample Input	116
20.4 Sample Output	116

Experiment 1: UNIX Permission and umask Calculator

Problem Statement

- Implement a CLI tool named permcalc.
- Inputs: --mode <octal> (required) and optional --umask <octal>.
- <octal> must be exactly 4 digits from 0000 to 0777 (leading zero required).
- Compute: effective_mode = mode & (~umask) (bitwise), limited to 0777.
- Output exactly two lines on success:
 - OK: EFFECTIVE <octal>
 - OK: SYMBOLIC <rwxrwxrwx>
- On any error, output exactly one line using the standard error format and exit non-zero.

Solution Implementation

```
1 import sys
2 import re
3
4 # Error codes
5 E_USAGE = "E_USAGE"
6 E_OCTAL = "E_OCTAL"
7 E_RANGE = "E_RANGE"
8
9 def error_exit(code, message):
10     """Print error message and exit with non-zero status."""
11     print(f"ERROR: {code}: {message}", file=sys.stderr)
12     sys.exit(1)
13
14 def parse_arguments(args):
```

```
15     """Parse command line arguments."""
16     mode = None
17     umask = "0000" # Default umask
18
19     i = 0
20
21     while i < len(args):
22         if args[i] == "--mode":
23             if i + 1 >= len(args):
24                 error_exit(E_USAGE, "missing value for --mode")
25             mode = args[i + 1]
26             i += 2
27
28         elif args[i] == "--umask":
29             if i + 1 >= len(args):
30                 error_exit(E_USAGE, "missing value for --umask")
31             umask = args[i + 1]
32             i += 2
33
34     else:
35         error_exit(E_USAGE, f"unrecognized argument: {args[i]}")
36
37     if mode is None:
38         error_exit(E_USAGE, "missing required --mode")
39
40     return mode, umask
41
42
43 def validate_octal(value, name):
44     """Validate that value is exactly 4 octal digits (0000-0777)."""
45
46     # Check length
47     if len(value) != 4:
48         return False, f"{name} must be exactly 4 digits"
49
50     # Check each character is octal digit
51     if not re.match(r'^[0-7]{4}$', value):
52         return False, f"{name} must be 4-digit octal (0000-0777)"
53
54     # Convert to integer and check range
55     try:
56         int_val = int(value, 8)
57         if int_val < 0 or int_val > 0o777:
```

```
53         return False, f"{name} must be in range 0000-0777"
54     except ValueError:
55         return False, f"{name} must be 4-digit octal (0000-0777)"
56
57     return True, None
58
59 def octal_to_symbolic(octal_str):
60     """Convert 4-digit octal string to symbolic rwxrwxrwx format."""
61     # We only care about last 3 digits (permissions)
62     # First digit is special permissions (sticky bit, setuid, setgid)
63     perm_str = octal_str[-3:] if len(octal_str) == 4 else octal_str
64
65     symbolic = ""
66     for digit in perm_str:
67         oct_val = int(digit)
68
69         # Build rwx for this digit
70         r = 'r' if oct_val & 4 else '-'
71         w = 'w' if oct_val & 2 else '-'
72         x = 'x' if oct_val & 1 else '-'
73
74         symbolic += r + w + x
75
76     return symbolic
77
78 def calculate_effective(mode_str, umask_str):
79     """Calculate effective mode = mode & (~umask)."""
80     mode_int = int(mode_str, 8)
81     umask_int = int(umask_str, 8)
82
83     # Apply umask: effective = mode AND (NOT umask)
84     # Mask with 0o777 to ensure we only get 9 permission bits
85     effective_int = mode_int & (~umask_int) & 0o777
86
87     # Format back to 4-digit octal with leading zeros
88     return f"{effective_int:04o}"
89
90 def main():
```

```
91 # Parse arguments (skip program name)
92 mode_str, umask_str = parse_arguments(sys.argv[1:])
93
94 # Validate inputs
95 valid, msg = validate_octal(mode_str, "mode")
96 if not valid:
97     error_exit(E_OCTAL, msg)
98
99 valid, msg = validate_octal(umask_str, "umask")
100 if not valid:
101     error_exit(E_OCTAL, msg)
102
103 # Calculate effective mode
104 effective_str = calculate_effective(mode_str, umask_str)
105
106 # Generate symbolic representation
107 symbolic_str = octal_to_symbolic(effective_str)
108
109 # Output results
110 print(f"OK: EFFECTIVE {effective_str}")
111 print(f"OK: SYMBOLIC {symbolic_str}")
112
113 sys.exit(0)
114
115 if __name__ == "__main__":
116     main()
```

Sample Input 1

```
$ python3 permcalc.py --mode 0644 --umask 0022
```

Sample Output 1

```
OK: EFFECTIVE 0644
OK: SYMBOLIC rw-r--r--
```

Sample Input 2

```
$ python3 permcalc.py --mode 0744 --umask 0022
```

Sample Output 2

```
OK: EFFECTIVE 0744
OK: SYMBOLIC rwxr--r--
```

Sample Input 3

```
$ python3 permcalc.py --mode 0844 --umask 0020
```

Sample Output 3

```
ERROR: E_OCTAL: mode must be 4-digit octal (0000-0777)
```

Experiment 2: POSIX File Copy with open/read/write

Problem Statement

- Implement a CLI tool named `fdcopy`.
- Inputs: `--src <path>` and `--dst <path>` (required), optional `--buf <N>` (1..1048576), optional `--force`.
- `--src` – means read from STDIN.
- Copy the exact byte stream from `src` to `dst` using only file-descriptor I/O.
- Compute CRC32 of bytes copied (IEEE 802.3) and total bytes copied.
- On success, output exactly two lines:
 - `OK: COPIED <bytes> BYTES`
 - `OK: CRC32 <8-hex>`
- If `dst` exists, fail unless `--force` is provided.
- On any error, output exactly one line using the standard error format and exit non-zero.

Solution Implementation

```
1  
2 import sys  
3 import os  
4 import zlib  
5  
6 # Error codes  
7 E_USAGE = "E_USAGE"  
8 E_OPEN_SRC = "E_OPEN_SRC"  
9 E_OPEN_DST = "E_OPEN_DST"  
10 E_EXISTS = "E_EXISTS"  
11 E_READ = "E_READ"
```

```
12 E_WRITE = "E_WRITE"
13 E_CLOSE = "E_CLOSE"
14 E_RANGE = "E_RANGE"
15
16 def error_exit(code, message):
17     """Print error message and exit with non-zero status."""
18     print(f"ERROR: {code}: {message}", file=sys.stderr)
19     sys.exit(1)
20
21 def parse_arguments(args):
22     """Parse command line arguments."""
23     src = None
24     dst = None
25     buffer_size = 4096 # Default buffer size
26     force = False
27
28     i = 0
29     while i < len(args):
30         if args[i] == "--src":
31             if i + 1 >= len(args):
32                 error_exit(E_USAGE, "missing value for --src")
33             src = args[i + 1]
34             i += 2
35         elif args[i] == "--dst":
36             if i + 1 >= len(args):
37                 error_exit(E_USAGE, "missing value for --dst")
38             dst = args[i + 1]
39             i += 2
40         elif args[i] == "--buf":
41             if i + 1 >= len(args):
42                 error_exit(E_USAGE, "missing value for --buf")
43             try:
44                 buffer_size = int(args[i + 1])
45             except ValueError:
46                 error_exit(E_USAGE, "buffer size must be an integer")
47             i += 2
48         elif args[i] == "--force":
49             force = True
```

```
50             i += 1
51
52         else:
53             error_exit(E_USAGE, f"unrecognized argument: {args[i]}")
54
55     if src is None:
56         error_exit(E_USAGE, "missing required --src")
57     if dst is None:
58         error_exit(E_USAGE, "missing required --dst")
59
60     # Validate buffer size
61     if buffer_size < 1 or buffer_size > 1048576: # 1..1048576 bytes
62         error_exit(E_RANGE, "buffer size must be 1..1048576 bytes")
63
64
65     return src, dst, buffer_size, force
66
67
68 def open_source(src_path):
69     """Open source file or stdin."""
70     if src_path == "-":
71         # Use stdin (file descriptor 0)
72         return 0 # STDIN_FILENO
73     else:
74         try:
75             # O_RDONLY: read only, no create
76             fd = os.open(src_path, os.O_RDONLY)
77             return fd
78         except OSError as e:
79             error_exit(E_OPEN_SRC, f"cannot open source: {e.strerror}")
80
81
82 def open_destination(dst_path, force):
83     """Open destination file with appropriate flags."""
84     # Default flags: write only, create, exclusive (fail if exists)
85     flags = os.O_WRONLY | os.O_CREAT
86
87     if not force:
88         flags |= os.O_EXCL # Fail if file exists
89
90
91     # Default permissions: rw-r--r-- (0644)
92     mode = 0o644
```

```
88
89     try:
90         fd = os.open(dst_path, flags, mode)
91         return fd
92     except OSError as e:
93         if e.errno == 17: # EEXIST: File exists
94             error_exit(E_EXISTS, "destination already exists (use --
95                         force)")
96         else:
97             error_exit(E_OPEN_DST, f"cannot open destination: {e.
98                         strerror}")
99
100
101 def copy_file(fd_src, fd_dst, buffer_size):
102     """
103     Copy from fd_src to fd_dst using buffer_size chunks.
104     Returns (bytes_copied, crc32_value)
105     """
106
107     bytes_copied = 0
108     crc32 = 0
109
110     while True:
111         # Read chunk from source
112         try:
113             chunk = os.read(fd_src, buffer_size)
114         except OSError as e:
115             error_exit(E_READ, f"read error: {e.strerror}")
116
117         # EOF reached
118         if not chunk:
119             break
120
121         # Write chunk to destination (handle partial writes)
122         bytes_written = 0
123         while bytes_written < len(chunk):
124             try:
125                 written = os.write(fd_dst, chunk[bytes_written:])
126             except OSError as e:
127                 error_exit(E_WRITE, f"write error: {e.strerror}")
```

```
124
125     if written == 0:
126         error_exit(E_WRITE, "write returned 0 bytes")
127
128     bytes_written += written
129
130     # Update CRC32 and byte count
131     crc32 = zlib.crc32(chunk, crc32)
132     bytes_copied += len(chunk)
133
134 return bytes_copied, crc32
135
136 def close_file(fd, fd_name):
137     """Close file descriptor with error handling."""
138     if fd > 0: # Don't close stdin (fd=0) if we didn't open it
139         try:
140             os.close(fd)
141         except OSError as e:
142             error_exit(E_CLOSE, f"cannot close {fd_name}: {e.strerror}")
143
144 def main():
145     # Parse arguments
146     src_path, dst_path, buffer_size, force = parse_arguments(sys.argv[1:])
147
148     # Open source
149     fd_src = open_source(src_path)
150
151     # Open destination
152     fd_dst = open_destination(dst_path, force)
153
154     try:
155         # Copy data
156         bytes_copied, crc32_value = copy_file(fd_src, fd_dst,
157             buffer_size)
158
159         # Ensure CRC32 is 32-bit unsigned
```

```
159     crc32_value = crc32_value & 0xffffffff
160
161     # Output results
162     print(f"OK: COPIED {bytes_copied} BYTES")
163     print(f"OK: CRC32 {crc32_value:08x}")
164
165 finally:
166     # Always close files
167     close_file(fd_src, "source")
168     close_file(fd_dst, "destination")
169
170 sys.exit(0)
171
172 if __name__ == "__main__":
173     main()
```

Sample Input 1

```
$ printf "abc" | python3 fdcopy.py --src - --dst out/test.bin
```

Sample Output 1

```
ERROR: E_OPEN_DST: cannot open destination: No such file or
directory
```

Sample Input 2

```
$ mkdir -p out
$ printf "abc" | python3 fdcopy.py --src - --dst out/test.bin
```

Sample Output 2

```
OK: COPIED 3 BYTES
```

```
OK: CRC32 352441c2
```

Experiment 3: Directory Listing and Metadata Report (ls + stat subset)

Problem Statement

- Implement a CLI tool named `dirreport`.
- Input: `--path <dir>` (required), optional `--sort name|size` (default: `name`).
- For each direct child entry (non-recursive), output one line: `ENTRY <type> <size> <name>`.
- `<type>` must be one of: F (regular file), D (directory), L (symlink), O (other).
- After listing, output a summary line: `OK: TOTAL <n> FILES <f> DIRS <d> LINKS <l> OTHER <o>`.
- If `--sort size`, sort by size ascending then name lexicographically.
- On error output exactly one `ERROR:` line and exit non-zero.

Solution Implementation

```
1
2 import sys
3 import os
4 import stat
5
6 # Error codes
7 E_USAGE = "E_USAGE"
8 E_NOTDIR = "E_NOTDIR"
9 E_OPEN_DIR = "E_OPEN_DIR"
10 E_READ_DIR = "E_READ_DIR"
11 E_STAT = "E_STAT"
12
13 def error_exit(code, message):
```

```
14     """Print error message and exit with non-zero status."""
15     print(f"ERROR: {code}: {message}", file=sys.stderr)
16     sys.exit(1)
17
18 def parse_arguments(args):
19     """Parse command line arguments."""
20     path = None
21     sort_by = "name" # default
22
23     i = 0
24     while i < len(args):
25         if args[i] == "--path":
26             if i + 1 >= len(args):
27                 error_exit(E_USAGE, "missing value for --path")
28             path = args[i + 1]
29             i += 2
30         elif args[i] == "--sort":
31             if i + 1 >= len(args):
32                 error_exit(E_USAGE, "missing value for --sort")
33             sort_by = args[i + 1]
34             if sort_by not in ["name", "size"]:
35                 error_exit(E_USAGE, "sort must be 'name' or 'size'")
36             i += 2
37         else:
38             error_exit(E_USAGE, f"unrecognized argument: {args[i]}")
39
40     if path is None:
41         error_exit(E_USAGE, "missing required --path")
42
43     return path, sort_by
44
45 def get_entry_type(mode):
46     """Determine entry type character from stat mode."""
47     if stat.S_ISREG(mode):
48         return 'F' # Regular file
49     elif stat.S_ISDIR(mode):
50         return 'D' # Directory
51     elif stat.S_ISLNK(mode):
```

```
52         return 'L' # Symbolic link
53     else:
54         return 'O' # Other (device, pipe, socket, etc.)
55
56 def list_directory(path, sort_by):
57     """
58     List directory entries and return list of (name, type, size) tuples
59     .
60     """
61
62     # Check if path exists and is a directory
63     try:
64         if not os.path.exists(path):
65             error_exit(E_NOTDIR, "path does not exist")
66
67         if not os.path.isdir(path):
68             error_exit(E_NOTDIR, "path is not a directory")
69     except OSError as e:
70         error_exit(E_NOTDIR, f"cannot access path: {e.strerror}")
71
72     # Open and read directory
73     try:
74         # Using scandir() which is more efficient than listdir()
75         with os.scandir(path) as it:
76             for entry in it:
77                 name = entry.name
78
79                 # Skip . and ..
80                 if name in ['.', '..']:
81                     continue
82
83                 # Get file info using lstat() (doesn't follow symlinks)
84                 try:
85                     stat_info = entry.stat(follow_symlinks=False)
86                 except OSError as e:
87                     error_exit(E_STAT, f"cannot stat '{name}': {e.strerror}")
```

```
88
89         # Determine type
90         entry_type = get_entry_type(stat_info.st_mode)
91
92         # Get size
93         size = stat_info.st_size
94
95         entries.append((name, entry_type, size))
96
97     except PermissionError as e:
98         error_exit(E_OPEN_DIR, f"permission denied: {e.strerror}")
99     except OSError as e:
100        error_exit(E_OPEN_DIR, f"cannot open directory: {e.strerror}")
101
102    # Sort entries
103    if sort_by == "name":
104        entries.sort(key=lambda x: x[0])    # Sort by name
105    elif sort_by == "size":
106        # Sort by size, then by name for ties
107        entries.sort(key=lambda x: (x[2], x[0]))
108
109    return entries
110
111 def main():
112     # Parse arguments
113     path, sort_by = parse_arguments(sys.argv[1:])
114
115     # Get directory entries
116     entries = list_directory(path, sort_by)
117
118     # Counters for summary
119     total = len(entries)
120     files = 0
121     dirs = 0
122     links = 0
123     other = 0
124
125     # Output each entry
```

```
126     for name, entry_type, size in entries:
127         print(f"ENTRY {entry_type} {size} {name}")
128
129         # Update counters
130         if entry_type == 'F':
131             files += 1
132         elif entry_type == 'D':
133             dirs += 1
134         elif entry_type == 'L':
135             links += 1
136         elif entry_type == 'O':
137             other += 1
138
139         # Output summary
140         print(f"OK: TOTAL {total} FILES {files} DIRS {dirs} LINKS {links}
141             OTHER {other}")
142
143
144 if __name__ == "__main__":
145     main()
```

Sample Input 1: Directory does not exist

```
$ python3 dirreport.py --path testdir --sort size
```

Sample Output 1

```
ERROR: E_NOTDIR: path does not exist
```

Commands to create test directory and files

```
cd LAB/
mkdir testdir
echo "Hello" > testdir/file1.txt
echo "World" > testdir/file2.txt
mkdir testdir/subdir
ln -s file1.txt testdir/link1
```

Sample Input 2: List directory sorted by size after creation

```
$ python3 dirreport.py --path testdir --sort size
```

Sample Output 2

```
ENTRY F 6 file1.txt
ENTRY F 6 file2.txt
ENTRY L 9 link1
ENTRY D 4096 subdir
OK: TOTAL 4 FILES 2 DIRS 1 LINKS 1 OTHER 0
```

Sample Input 3: Directory **anotherdir**

```
$ mkdir anotherdir
$ echo "Data" > anotherdir/data1.txt
$ echo "MoreData" > anotherdir/data2.txt
$ mkdir anotherdir/subfolder
$ ln -s data1.txt anotherdir/linkA
$ python3 dirreport.py --path anotherdir --sort size
```

Sample Output 3

```
ENTRY F 5 data1.txt
ENTRY F 9 data2.txt
ENTRY L 6 linkA
ENTRY D 4096 subfolder
OK: TOTAL 4 FILES 2 DIRS 1 LINKS 1 OTHER 0
```

Sample Input 4: Directory **missingdir** (does not exist)

```
$ python3 dirreport.py --path missingdir
```

Sample Output 4

```
ERROR: E_NOTDIR: path does not exist
```

Experiment 4: grep-lite: Deterministic Text Pattern Search

Problem Statement

- Implement a CLI tool named greplite.
- Inputs: --pattern <ASCII> (required), --files <f1, f2, ...> (required).
- Match is literal substring (case-sensitive) within each line.
- For each match, output a line: MATCH <file>:<line_no>:<line> where <line> is the original line without trailing newline.
- After processing all files, output exactly one summary line: OK: MATCHES <k> FILES <n> (<k> = total matches, <n> = number of files processed).
- If any file cannot be opened, treat it as an error (do not partially succeed).
- On any error, output exactly one ERROR: line and exit non-zero.

Solution Implementation

```
1 import sys
2
3 # Error codes
4 E_USAGE = "E_USAGE"
5 E_EMPTY_PATTERN = "E_EMPTY_PATTERN"
6 E_OPEN = "E_OPEN"
7 E_READ = "E_READ"
8
9 def error_exit(code, message):
10     """Print error message and exit with non-zero status."""
11     print(f"ERROR: {code}: {message}", file=sys.stderr)
12     sys.exit(1)
13
```

```
14 def parse_arguments(args):
15     """Parse command line arguments."""
16     pattern = None
17     files_str = None
18
19     i = 0
20     while i < len(args):
21         if args[i] == "--pattern":
22             if i + 1 >= len(args):
23                 error_exit(E_USAGE, "missing value for --pattern")
24             pattern = args[i + 1]
25             i += 2
26         elif args[i] == "--files":
27             if i + 1 >= len(args):
28                 error_exit(E_USAGE, "missing value for --files")
29             files_str = args[i + 1]
30             i += 2
31         else:
32             error_exit(E_USAGE, f"unrecognized argument: {args[i]}")
33
34     if pattern is None:
35         error_exit(E_USAGE, "missing required --pattern")
36     if files_str is None:
37         error_exit(E_USAGE, "missing required --files")
38
39     # Validate pattern
40     if pattern == "":
41         error_exit(E_EMPTY_PATTERN, "pattern must be non-empty")
42
43     # Parse file list
44     if files_str.endswith(','):
45         error_exit(E_USAGE, "invalid file list: trailing comma")
46
47     files = [f.strip() for f in files_str.split(',') if f.strip()]
48
49     if not files:
50         error_exit(E_USAGE, "must provide at least one file")
51
```

```
52     # Check for empty filenames in the middle
53     if '' in [f.strip() for f in files_str.split(',')]:
54         error_exit(E_USAGE, "file list contains empty entries")
55
56     return pattern, files
57
58 def search_files(pattern, files):
59     """Search for pattern in files and return results."""
60     matches = []
61     total_matches = 0
62     files_processed = 0
63
64     # Try to open all files first (all-or-nothing)
65     file_handles = []
66     try:
67         for filename in files:
68             try:
69                 f = open(filename, 'r')
70                 file_handles.append((filename, f))
71             except OSError as e:
72                 # Close any files we successfully opened
73                 for _, f in file_handles:
74                     f.close()
75                 error_exit(E_OPEN, f"cannot open '{filename}': {e.strerror}")
76             except Exception as e:
77                 # Catch any other unexpected errors
78                 for _, f in file_handles:
79                     f.close()
80                 error_exit(E_OPEN, f"unexpected error opening files: {e}")
81
82     # Now search each file
83     for filename, f in file_handles:
84         try:
85             line_number = 0
86             file_match_count = 0
87
88             for line in f:
```

```
89         line_number += 1
90         # Remove trailing newline but keep other whitespace
91         line_content = line.rstrip('\n')
92
93         # Check for pattern (case-sensitive substring)
94         if pattern in line_content:
95             matches.append(f"MATCH {filename}:{line_number}:{{
96                 line_content}}")
97             total_matches += 1
98             file_match_count += 1
99
100            files_processed += 1
101
102        except OSError as e:
103            # Close all files before exiting
104            for _, f in file_handles:
105                f.close()
106            error_exit(E_READ, f"error reading '{filename}': {e.
107                         strerror}")
108
109        finally:
110            f.close()
111
112    return matches, total_matches, files_processed
113
114
115 def main():
116     # Parse arguments
117     pattern, files = parse_arguments(sys.argv[1:])
118
119     # Search files
120     matches, total_matches, files_processed = search_files(pattern,
121                 files)
122
123     # Output matches
124     for match in matches:
125         print(match)
126
127
128     # Output summary
129     print(f"OK: MATCHES {total_matches} FILES {files_processed}")
```

```
124  
125     sys.exit(0)  
126  
127 if __name__ == "__main__":  
128     main()
```

Commands to create test files

```
cd "LAB"  
echo "abc def" > test1.txt  
echo "hello" > test2.txt
```

Sample Input 1: Files not yet created

```
$ python3 4.greplite.py --pattern "xyz" --files "test1.txt,  
test2.txt"
```

Sample Output 1: Files missing

```
ERROR: E_OPEN: cannot open 'test1.txt': No such file or  
directory
```

Sample Input 2: After creating test files

```
$ python3 4.greplite.py --pattern "xyz" --files "test1.txt,  
test2.txt"
```

Sample Output 2: No matches

```
OK: MATCHES 0 FILES 2
```

Sample Input 3: Pattern "hello"

```
$ python3 4.greplite.py --pattern "hello" --files "test1.txt,  
test2.txt"
```

Sample Output 3: One match

```
MATCH test2.txt:1:hello  
OK: MATCHES 1 FILES 2
```

Experiment 5: Process Spawner and Exit-Status Reporter (fork/exec/wait)

Problem Statement

- Implement a CLI tool named spawnwait.
- Inputs: --cmd <program> (required), optional --args <a₁, a₂, ...> and optional --repeat <k> (default 1).
- Spawn k children sequentially (next starts after previous terminates).
- For each child, print:
 - CHILD <i> PID <pid> START
 - Normal exit: CHILD <i> PID <pid> EXIT <code>
 - Signal termination: CHILD <i> PID <pid> SIG <signum>
- After all children complete, print: OK: COMPLETED <k>.
- On any error (fork/exec/wait failures), output one ERROR: line and exit non-zero.

Solution Implementation

```
1 import sys
2 import os
3
4 def error_exit(code, message):
5     print(f"ERROR: {code}: {message}", file=sys.stderr)
6     sys.exit(1)
7
8 def main():
9     # 1. Parse Arguments
10    args = sys.argv[1:]
11    cmd = None
```

```
12     cmd_args = []
13     repeat = 1
14
15     i = 0
16     while i < len(args):
17         if args[i] == "--cmd":
18             if i + 1 < len(args):
19                 cmd = args[i+1]
20                 i += 2
21             else:
22                 error_exit("E_USAGE", "missing value for --cmd")
23         elif args[i] == "--args":
24             if i + 1 < len(args):
25                 # Split comma-separated arguments
26                 cmd_args = args[i+1].split(',')
27                 i += 2
28             else:
29                 error_exit("E_USAGE", "missing value for --args")
30         elif args[i] == "--repeat":
31             if i + 1 < len(args):
32                 try:
33                     repeat = int(args[i+1])
34                     if repeat < 1: raise ValueError
35                 except ValueError:
36                     error_exit("E_RANGE", "repeat must be >= 1")
37                     i += 2
38                 else:
39                     error_exit("E_USAGE", "missing value for --repeat")
40             else:
41                 error_exit("E_USAGE", f"unrecognized argument: {args[i]}")
42
43     if cmd is None:
44         error_exit("E_USAGE", "missing required --cmd")
45
46     # 2. Sequential Spawning Loop
47     for k in range(1, repeat + 1):
48         try:
49             # Step 1: Fork
```

```
50         pid = os.fork()
51     except OSError:
52         error_exit("E_FORK", "failed to fork process")
53
54     if pid == 0:
55         # --- CHILD PROCESS ---
56         try:
57             # Step 2: Exec
58             # os.execvp(file, args_list) - first arg must be the
59             # program name
60             os.execvp(cmd, [cmd] + cmd_args)
61         except OSError:
62             # If exec fails, the child must exit immediately
63             # We exit with a special code so parent knows it's an
64             # exec failure
65             os._exit(127)
66
67     else:
68         # --- PARENT PROCESS ---
69         print(f"CHILD {k} PID {pid} START")
70
71         # Step 3: Wait for termination
72         try:
73             # waitpid(pid, options)
74             _, status = os.waitpid(pid, 0)
75
76             # Interpret status
77             if os.WIFEXITED(status):
78                 exit_code = os.WEXITSTATUS(status)
79                 # Check if the exit was actually an exec failure
80                 if exit_code == 127:
81                     error_exit("E_EXEC", "cannot exec program")
82                 print(f"CHILD {k} PID {pid} EXIT {exit_code}")
83
84             elif os.WIFSIGNALED(status):
85                 signum = os.WTERMSIG(status)
86                 print(f"CHILD {k} PID {pid} SIG {signum}")
87
88         except OSError:
```

```
86             error_exit("E_WAIT", "waitpid failed")
87
88     print(f"OK: COMPLETED {repeat}")
89
90 if __name__ == "__main__":
91     main()
```

Sample Input 1

```
$ python3 spawnwait.py --cmd /bin/true --repeat 3
```

Sample Output 1

```
CHILD 1 PID 34879 START
CHILD 1 PID 34879 EXIT 0
CHILD 2 PID 34880 START
CHILD 2 PID 34880 EXIT 0
CHILD 3 PID 34881 START
CHILD 3 PID 34881 EXIT 0
OK: COMPLETED 3
```

Sample Input 2: Invalid command

```
$ python3 spawnwait.py --cmd /bin/sh --args -c,exit\ 7
```

Sample Output 2

```
CHILD 1 PID 34941 START
CHILD 1 PID 34941 EXIT 7
OK: COMPLETED 1
```

Experiment 6: Signal-Based Timeout Supervisor (sigaction + alarm + kill)

Problem Statement

- Implement a CLI tool named `timeoutwrap`.
- Inputs: `--seconds <t>` (required, integer 1..60), `--cmd <program>` (required), optional `--args <a1, a2, ...>`.
- The parent process must fork a child that executes the given command.
- The parent arms an alarm for t seconds.
- If the child exits before timeout, cancel the alarm and output: `OK: EXIT <code>`.
- If timeout occurs first, send `SIGKILL` to the child, wait for it, and output: `OK: TIMEOUT KILLED`.
- On any error (fork, exec, signal, wait), output exactly one `ERROR: line` and exit non-zero.

Solution Implementation

```
1 import sys
2 import os
3 import signal
4 import time
5
6 # Error codes
7 E_USAGE = "E_USAGE"
8 E_RANGE = "E_RANGE"
9 E_FORK = "E_FORK"
10 E_EXEC = "E_EXEC"
11 E_WAIT = "E_WAIT"
12 E_SIGNAL = "E_SIGNAL"
```

```
13
14 def error_exit(code, message):
15     print(f"ERROR: {code}: {message}", file=sys.stderr)
16     sys.exit(1)
17
18 # Global variable to track the child PID
19 child_pid = -1
20
21 def alarm_handler(signum, frame):
22     """This function runs when the alarm timer reaches zero."""
23     global child_pid
24     if child_pid > 0:
25         try:
26             # Send SIGKILL to the child
27             os.kill(child_pid, signal.SIGKILL)
28         except OSError:
29             pass # Child might have just finished
30
31 def main():
32     global child_pid
33
34     # 1. Parse Arguments
35     args = sys.argv[1:]
36     seconds = None
37     cmd = None
38     cmd_args = []
39
40     i = 0
41     while i < len(args):
42         if args[i] == "--seconds":
43             if i + 1 < len(args):
44                 try:
45                     seconds = int(args[i+1])
46                     if not (1 <= seconds <= 60): raise ValueError
47                 except ValueError:
48                     error_exit(E_RANGE, "seconds must be in 1..60")
49             i += 2
50         else:
```

```
51             error_exit(E_USAGE, "missing value for --seconds")
52     elif args[i] == "--cmd":
53         if i + 1 < len(args):
54             cmd = args[i+1]
55             i += 2
56     else:
57         error_exit(E_USAGE, "missing value for --cmd")
58     elif args[i] == "--args":
59         if i + 1 < len(args):
60             cmd_args = args[i+1].split(',')
61             i += 2
62         else:
63             error_exit(E_USAGE, "missing value for --args")
64     else:
65         error_exit(E_USAGE, f"unrecognized argument: {args[i]}")
66
67     if seconds is None or cmd is None:
68         error_exit(E_USAGE, "--seconds and --cmd are required")
69
70     # 2. Setup Signal Handler
71     # signal.signal is the Python equivalent to sigaction()
72     signal.signal(signal.SIGALRM, alarm_handler)
73
74     # 3. Fork and Exec
75     try:
76         child_pid = os.fork()
77     except OSError:
78         error_exit(E_FORK, "fork failed")
79
80     if child_pid == 0:
81         # --- CHILD PROCESS ---
82         try:
83             os.execvp(cmd, [cmd] + cmd_args)
84         except OSError:
85             os._exit(127) # Exec failure
86     else:
87         # --- PARENT PROCESS ---
88         # Arm the alarm
```

```
89     signal.alarm(seconds)

90

91     try:
92         # Wait for the child to change state
93         pid, status = os.waitpid(child_pid, 0)
94
95         # Cancel the alarm because the child finished
96         signal.alarm(0)
97
98         # Analyze how the child died
99         if os.WIFEXITED(status):
100             code = os.WEXITSTATUS(status)
101             if code == 127:
102                 error_exit(E_EXEC, "cannot exec program")
103             print(f"OK: EXIT {code}")
104
105         elif os.WIFSIGNALED(status):
106             sig = os.WTERMSIG(status)
107             if sig == signal.SIGKILL:
108                 # This means our alarm_handler killed it
109                 print("OK: TIMEOUT KILLED")
110             else:
111                 # Child died by some other signal (e.g. SIGINT)
112                 print(f"OK: SIG {sig}")
113
114     except OSError:
115         error_exit(E_WAIT, "waitpid failed")
116
117 if __name__ == "__main__":
118     main()
```

Sample Input 1: Command completes before timeout

```
$ python3 timeoutwrap.py --seconds 2 --cmd /bin/sh --args -c,
sleep\ 1
```

Sample Output 1

```
OK: EXIT 0
```

Sample Input 2: Command exceeds timeout

```
$ python3 timeoutwrap.py --seconds 1 --cmd /bin/sh --args -c,  
sleep\ 7
```

Sample Output 2

```
OK: TIMEOUT KILLED
```

Sample Input 3

```
$ python3 timeoutwrap.py --seconds 70 --cmd /bin/true
```

Sample Output 3

```
ERROR: E_RANGE: seconds must be in 1..60
```

Experiment 7: Pipe-Based Filter Chain (pipe + dup2)

Problem Statement

- Implement a CLI tool named pipechain.
- Inputs: --producer <cmd1>, --filter <cmd2>, --consumer <cmd3> (all required).
- Each <cmd> is a single shell-free command path with optional comma-separated arguments via --producer-args, --filter-args, --consumer-args.
- Run the equivalent of cmd1 | cmd2 | cmd3 using two pipes and three child processes.
- Parent waits for all children; if all exit 0, print: OK: PIPELINE SUCCESS.
- If any stage exits non-zero, print exactly one line: ERROR: E_STAGE: stage <name> exit <code>.
- If a stage is terminated by a signal, print exactly one line: ERROR: E_STAGE: stage <name> sig <signum>.

Solution Implementation

```
1 import sys
2 import os
3
4 def error_exit(msg):
5     print(f"ERROR: {msg}")
6     sys.exit(1)
7
8 def main():
9     # 1. Argument Parsing
10    args_raw = sys.argv[1:]
11    params = {}
12    i = 0
```

```
13     while i < len(args_raw):
14         key = args_raw[i].lstrip('-')
15         if i + 1 < len(args_raw):
16             params[key] = args_raw[i+1]
17             i += 2
18         else:
19             error_exit("E_USAGE: Missing value for argument")
20
21     # Required check
22     for req in ['producer', 'filter', 'consumer']:
23         if req not in params:
24             error_exit("E_USAGE: Missing required stage")
25
26     # Prepare command lists
27     stages = [
28         ("producer", params['producer'], params.get('producer-args', '').
29          split(',') if params.get('producer-args') else []),
30         ("filter", params['filter'], params.get('filter-args', '').
31          split(',') if params.get('filter-args') else []),
32         ("consumer", params['consumer'], params.get('consumer-args', '').
33          split(',') if params.get('consumer-args') else [])
34     ]
35
36     # 2. Create Pipes
37     # pipel: Producer -> Filter | pipe2: Filter -> Consumer
38     p1_read, p1_write = os.pipe()
39     p2_read, p2_write = os.pipe()
40
41     pids = {}
42
43     # 3. Spawn Stages
44     for name, cmd, cmd_args in stages:
45         try:
46             pid = os.fork()
47             if pid == 0: # CHILD
48                 # Redirect STDERR and STDOUT to /dev/null as per spec
49                 devnull = os.open(os.devnull, os.O_WRONLY)
50                 os.dup2(devnull, sys.stderr.fileno())
51
52                 # Create pipes
53                 p1, p2 = os.pipe()
54
55                 # Set up producer
56                 if name == "producer":
57                     os.setfd(1, p2)
58
59                 # Set up filter
60                 if name == "filter":
61                     os.setfd(0, p1)
62                     os.setfd(1, p2)
63
64                 # Set up consumer
65                 if name == "consumer":
66                     os.setfd(0, p1)
```

```
48
49         if name == "producer":
50             os.dup2(p1_write, sys.stdout.fileno())
51         elif name == "filter":
52             os.dup2(p1_read, sys.stdin.fileno())
53             os.dup2(p2_write, sys.stdout.fileno())
54         elif name == "consumer":
55             os.dup2(p2_read, sys.stdin.fileno())
56             os.dup2(devnull, sys.stdout.fileno()) # Consumer
57                                         output to devnull
58
59         # Close all pipe fds in child
60         for fd in [p1_read, p1_write, p2_read, p2_write,
61                     devnull]:
62             os.close(fd)
63
64
65         pids[name] = pid
66     except OSError:
67         error_exit("E_FORK: Fork failed")
68
69     # 4. PARENT: Close all pipe ends
70     for fd in [p1_read, p1_write, p2_read, p2_write]:
71         os.close(fd)
72
73     # 5. Wait and Report (Check in fixed order: producer, filter,
74     # consumer)
75     final_error = None
76     for name, _, _ in stages:
77         _, status = os.waitpid(pids[name], 0)
78
79         if final_error is None: # Only capture the first error found in
80                               # order
81             if os.WIFEXITED(status):
82                 code = os.WEXITSTATUS(status)
83                 if code != 0:
```

```
82             # code 127 is our custom exec failure
83             msg = "cannot exec" if code == 127 else f"exit {code}"
84             final_error = f"E_STAGE: stage {name} {msg}"
85         elif os.WIFSIGNALED(status):
86             signum = os.WTERMSIG(status)
87             final_error = f"E_STAGE: stage {name} sig {signum}"
88
89     if final_error:
90         print(f"ERROR: {final_error}")
91         sys.exit(1)
92     else:
93         print("OK: PIPELINE SUCCESS")
94
95 if __name__ == "__main__":
96     main()
```

Sample Input 1

```
$ python3 pipechain.py --producer /bin/echo --producer-args
world --filter /usr/bin/tr --filter-args a-z,A-Z --consumer
/usr/bin/wc --consumer-args -c
```

Sample Output 1

```
OK: PIPELINE SUCCESS
```

Sample Input 2: Producer fails

```
$ python3 pipechain.py --producer /bin/sh --producer-args  
-c,exit\ 4 --filter /bin/cat --consumer /bin/true
```

Sample Output 2

```
ERROR: E_STAGE: stage producer exit 4
```

Experiment 8: Shared Memory Counter IPC (shm_open + mmap + sem_open)

Problem Statement

- Implement a CLI tool named `shmcounter`.
- Inputs: `--procs <p>` (2..16), `--iters <n>` (1..100000), `--name <id>` (alphanumeric, 1..16).
- Create shared memory object `/shm_<id>` containing a 64-bit signed integer counter initialized to 0.
- Create named semaphore `/sem_<id>` initialized to 1.
- Fork `p` child processes; each performs `n` increments of the shared counter with semaphore protection.
- After all children exit, output exactly: `OK: FINAL <value>` where `<value>=p*n`.
- Always unlink shared memory and semaphore before exit (success or failure).
- On error, output one `ERROR:` line and exit non-zero.

Solution Implementation

```
1 import sys
2 import os
3 import argparse
4 import mmap
5 import struct
6 import posix_ipc # Note: 'pip install posix_ipc' is usually required
    for POSIX semaphores
7 import time
8
9 # Error codes
```

```
10 E_USAGE = "E_USAGE"
11 E_RANGE = "E_RANGE"
12 E_SHM = "E_SHM"
13 E_MMAP = "E_MMAP"
14 E_SEM = "E_SEM"
15 E_FORK = "E_FORK"
16 E_WAIT = "E_WAIT"

17
18 def error_exit(code, message):
19     print(f"ERROR: {code}: {message}")
20     sys.exit(1)

21
22 def main():
23     # 1. Parse and Validate Arguments
24     parser = argparse.ArgumentParser(add_help=False)
25     parser.add_argument('--procs', type=int)
26     parser.add_argument('--iters', type=int)
27     parser.add_argument('--name')
28     args, _ = parser.parse_known_args()

29
30     if args.procs is None or args.iters is None or args.name is None:
31         error_exit(E_USAGE, "missing required arguments")

32
33     if not (2 <= args.procs <= 16):
34         error_exit(E_RANGE, "procs must be in 2..16")
35     if not (1 <= args.iters <= 100000):
36         error_exit(E_RANGE, "iters must be in 1..100000")
37     if not args.name.isalnum():
38         error_exit(E_RANGE, "name must be alphanumeric only")

39
40     shm_name = f"/shm_{args.name}"
41     sem_name = f"/sem_{args.name}"

42
43     shm = None
44     sem = None

45
46     try:
47         # 2. Create and Map Shared Memory
```

```
48     try:
49         # Create a 8-byte shared memory object (for 64-bit int)
50         shm = posix_ipc.SharedMemory(shm_name, flags=posix_ipc.
51             O_CREAT | posix_ipc.O_TRUNC, size=8)
52         map_file = mmap.mmap(shm.fd, 8)
53         # Initialize counter to 0
54         map_file.seek(0)
55         map_file.write(struct.pack('q', 0))
56     except Exception as e:
57         error_exit(E_SHM, f"could not create shm: {e}")
58
59     # 3. Create Semaphore
60     try:
61         sem = posix_ipc.Semaphore(sem_name, flags=posix_ipc.O_CREAT
62             | posix_ipc.O_TRUNC, initial_value=1)
63     except Exception as e:
64         error_exit(E_SEM, f"could not create semaphore: {e}")
65
66     # 4. Fork Processes
67     pids = []
68     for i in range(args.procs):
69         try:
70             pid = os.fork()
71             if pid == 0: # CHILD
72                 child_shm = posix_ipc.SharedMemory(shm_name)
73                 child_map = mmap.mmap(child_shm.fd, 8)
74                 child_sem = posix_ipc.Semaphore(sem_name)
75
76                 for _ in range(args.iters):
77                     child_sem.acquire()
78                     # Read 64-bit int, increment, and write back
79                     child_map.seek(0)
80                     val = struct.unpack('q', child_map.read(8))[0]
81                     child_map.seek(0)
82                     child_map.write(struct.pack('q', val + 1))
83                     child_sem.release()
84
85                 child_map.close()
86
87             else:
```

```
84                 os._exit(0)
85
86             else:
87
88                 pids.append(pid)
89
90         except OSError:
91
92             error_exit(E_FORK, "fork failed")
93
94
95     # 5. Parent Wait
96
97     for pid in pids:
98
99         try:
100             os.waitpid(pid, 0)
101
102         except OSError:
103             error_exit(E_WAIT, "waitpid failed")
104
105
106     # 6. Read Final Value
107
108     map_file.seek(0)
109
110     final_val = struct.unpack('q', map_file.read(8))[0]
111
112     print(f"OK: FINAL {final_val}")
113
114
115     finally:
116
117         # 7. Cleanup (Unlink)
118
119         if shm:
120
121             shm.close_fd()
122
123             try: posix_ipc.unlink_shared_memory(shm_name)
124
125             except: pass
126
127         if sem:
128
129             sem.close()
130
131             try: posix_ipc.unlink_semaphore(sem_name)
132
133             except: pass
134
135
136 if __name__ == "__main__":
137
138     main()
```

Sample Input 1

```
$ python3 shmcounter.py --procs 16 --iters 100000 --name t
```

Sample Output 1

```
OK: FINAL 1600000
```

Sample Input 2: Invalid Arguments

```
$ python3 shmcounter.py --procs 2 --iters 0 --name t
```

Sample Output 2

```
ERROR: E_RANGE: iters must be in 1..100000
```

Experiment 9: Threaded Deterministic Reducer (pthreads + mutex)

Problem Statement

- Implement a CLI tool named `thrsum`.
- Inputs: `--threads <t>` (1..32) and `--n <N>` (1..1000000).
- Compute the sum of integers 1..N using t threads.
- Work partition must be deterministic: thread *i* handles a contiguous block of the range.
- Each thread computes a local sum and then adds to a shared total under a mutex.
- Output exactly one line: `OK: SUM <value>`.
- On error, output one `ERROR: line` and exit non-zero.

Solution Implementation

```
1 import threading
2 import sys
3
4 def error(code, message):
5     print(f"ERROR: {code}: {message}")
6     sys.exit(1)
7
8 def worker(start, end, total, lock):
9     local_sum = 0
10    for i in range(start, end + 1):
11        local_sum += i
12
13    # protect shared total
14    with lock:
15        total[0] += local_sum
```

```
16
17 def main():
18     # very simple argument parsing
19     if "--threads" not in sys.argv or "--n" not in sys.argv:
20         error("E_USAGE", "missing required arguments")
21
22     try:
23         t = int(sys.argv[sys.argv.index("--threads") + 1])
24         n = int(sys.argv[sys.argv.index("--n") + 1])
25     except (ValueError, IndexError):
26         error("E_USAGE", "invalid arguments")
27
28     if t < 1 or t > 32:
29         error("E_RANGE", "threads must be in 1..32")
30
31     if n < 1 or n > 1_000_000:
32         error("E_RANGE", "n must be in 1..1000000")
33
34     threads = []
35     lock = threading.Lock()
36     total = [0] # mutable container for shared sum
37
38     chunk = n // t
39     remainder = n % t
40     start = 1
41
42     for i in range(t):
43         end = start + chunk - 1
44         if i < remainder:
45             end += 1
46
47         th = threading.Thread(
48             target=worker,
49             args=(start, end, total, lock)
50         )
51         threads.append(th)
52         th.start()
53
```

```
54         start = end + 1
55
56     for th in threads:
57         th.join()
58
59     print(f"OK: SUM {total[0]}")
60
61 if __name__ == "__main__":
62     main()
```

Sample Input 1

```
$ $ python3 threadsum.py --threads 4 --n 20
```

Sample Output 1

```
OK: SUM 210
```

Sample Input 2: Invalid Arguments

```
$ ./thrsum --threads 50 --n 500
```

Sample Output 2

```
ERROR: E_RANGE: threads must be in 1..32
```

Experiment 10: Bounded Buffer Producer-Consumer with Semaphores

Problem Statement

- Implement a CLI tool named pcbuf.
- Inputs: --buf (1..1024), --producers <p> (1..16), --consumers <c> (1..16), --items <m> (1..100000).
- Total items produced must equal m and total items consumed must equal m.
- Each produced item is the integer sequence 1..m (assigned in increasing order by a protected counter).
- Consumers compute the sum of consumed values; after all threads join, output exactly:

```
OK: PRODUCED <m>
OK: CONSUMED <m>
OK: SUM <s> # where S = m*(m+1)/2
```

- On error, output one ERROR: line and exit non-zero.

Solution Implementation

```
1 #!/usr/bin/env python3
2 import sys
3 import argparse
4 import threading
5 from queue import Queue
6 from threading import Semaphore, Lock
7
8 def parse_args():
9     parser = argparse.ArgumentParser()
```

```
10    parser.add_argument("--buf", type=int, required=True)
11    parser.add_argument("--producers", type=int, required=True)
12    parser.add_argument("--consumers", type=int, required=True)
13    parser.add_argument("--items", type=int, required=True)
14    return parser.parse_args()
15
16 def main():
17     try:
18         args = parse_args()
19         B = args.buf
20         P = args.producers
21         C = args.consumers
22         M = args.items
23
24         if not (1 <= B <= 1024 and 1 <= P <= 16 and 1 <= C <= 16 and 1
25             <= M <= 100000):
26             print("ERROR: Invalid arguments")
27             sys.exit(1)
28
29         buffer = Queue(maxsize=B)
30         empty_slots = Semaphore(B)
31         filled_slots = Semaphore(0)
32         counter_lock = Lock()
33         consume_lock = Lock()
34
35         produced_count = [0]
36         consumed_count = [0]
37         consumed_sum = [0]
38         next_item = [1] # shared counter for items 1..M
39
40         def producer():
41             nonlocal next_item
42             while True:
43                 with counter_lock:
44                     if next_item[0] > M:
45                         break
46                     item = next_item[0]
47                     next_item[0] += 1
48
49         producer()
50
51         for i in range(M):
52             with consumer_lock:
53                 empty_slots.acquire()
54                 buffer.put(i)
55                 filled_slots.release()
56
57         for i in range(C):
58             with consumer_lock:
59                 filled_slots.acquire()
60                 item = buffer.get()
61                 consumed_count[0] += 1
62                 consumed_sum[0] += item
63                 if consumed_count[0] == M:
64                     break
65
66         print(f"Consumed {consumed_count[0]} items with sum {consumed_sum[0]}")
67
68     except KeyboardInterrupt:
69         print("Program interrupted by user")
70
71     finally:
72         empty_slots.close()
73         filled_slots.close()
74         counter_lock.close()
75         consume_lock.close()
```

```
47         empty_slots.acquire()
48         buffer.put(item)
49         filled_slots.release()
50         with counter_lock:
51             produced_count[0] += 1
52
53     def consumer():
54         while True:
55             filled_slots.acquire()
56             if consumed_count[0] >= M:
57                 filled_slots.release()
58                 break
59             item = buffer.get()
60             empty_slots.release()
61             with consume_lock:
62                 consumed_count[0] += 1
63                 consumed_sum[0] += item
64
65     producers = [threading.Thread(target=producer) for _ in range(P)]
66     consumers = [threading.Thread(target=consumer) for _ in range(C)]
67
68     for th in producers + consumers:
69         th.start()
70     for th in producers + consumers:
71         th.join()
72
73     print(f"OK: PRODUCED {produced_count[0]}")
74     print(f"OK: CONSUMED {consumed_count[0]}")
75     print(f"OK: SUM {consumed_sum[0]}")
76
77     except Exception as e:
78         print("ERROR:", e)
79         sys.exit(1)
80
81 if __name__ == "__main__":
82     main()
```

Sample Input 1

```
$ python3 pcbuf.py --buf 10 --producers 2 --consumers 2  
--items 100
```

Sample Output 1

```
OK: PRODUCED 100  
OK: CONSUMED 100  
OK: SUM 5050
```

Sample Input 2: Invalid Arguments

```
$ python3 pcbuf.py --buf 0 --producers 2 --consumers 2  
--items 50
```

Sample Output 2

```
ERROR: E_RANGE: buf must be in 1..1024
```

Experiment 11: CPU Scheduling Simulator I (FCFS and Non-preemptive SJF)

Problem Statement

- Implement a CLI tool named `schedsim1`.
- Input is provided via `stdin` as CSV with header: `pid,arrival,burst` (`pid` is a string without commas).
- Simulate FCFS and non-preemptive SJF:
 - FCFS: First-Come, First-Served.
 - SJF: Non-preemptive Shortest Job First; among arrived processes, choose the shortest burst, tie-break by arrival then pid.
- For each algorithm, output exactly:

```
ALG <name>
GANTT <pid1>@<t0>-<t1> <pid2>@<t1>-<t2> ...
OK: AVG_WAIT <w> AVG_TAT <t>
```

Averages are rounded to 2 decimal places.

- On any parse or validation error, output one `ERROR:` line and exit non-zero.

Solution Implementation

```
1
2 import sys
3 from collections import namedtuple
4
5 Process = namedtuple("Process", ["pid", "arrival", "burst"])
6
7 def parse_input():
```

```
8     lines = [line.strip() for line in sys.stdin if line.strip()]
9     if not lines:
10         print("ERROR: E_INPUT: empty input", file=sys.stderr)
11         sys.exit(1)
12     if lines[0] != "pid,arrival,burst":
13         print("ERROR: E_INPUT: missing or incorrect header", file=sys.
14             stderr)
15         sys.exit(1)
16
16     processes = []
17     seen_pids = set()
18     for i, line in enumerate(lines[1:], start=2):
19         parts = line.split(",")
20         if len(parts) != 3:
21             print(f"ERROR: E_INPUT: line {i} malformed", file=sys.
22                 stderr)
22             sys.exit(1)
23         pid = parts[0].strip()
24         if pid in seen_pids:
25             print(f"ERROR: E_DUPPID: duplicate pid {pid}", file=sys.
26                 stderr)
26             sys.exit(1)
27         seen_pids.add(pid)
28         try:
29             arrival = int(parts[1].strip())
30             burst = int(parts[2].strip())
31         except ValueError:
32             print(f"ERROR: E_INPUT: arrival and burst must be integers
33                 on line {i}", file=sys.stderr)
33             sys.exit(1)
34         if arrival < 0 or burst <= 0:
35             print("ERROR: E_RANGE: arrival and burst must be non-
36                 negative; burst must be > 0", file=sys.stderr)
36             sys.exit(1)
37         processes.append(Process(pid, arrival, burst))
38     return processes
39
40 def fcfs(processes):
```

```
41     processes.sort(key=lambda p: p.arrival)
42     time = 0
43     gantt = []
44     wait_times = []
45     tat_times = []
46
47     for p in processes:
48         if time < p.arrival:
49             gantt.append(f"IDLE@{time}-{p.arrival}")
50             time = p.arrival
51         start = time
52         end = start + p.burst
53         gantt.append(f"{p.pid}@{start}-{end}")
54         wait_times.append(start - p.arrival)
55         tat_times.append(end - p.arrival)
56         time = end
57
58     avg_wait = round(sum(wait_times)/len(wait_times), 2)
59     avg_tat = round(sum(tat_times)/len(tat_times), 2)
60
61     print("ALG FCFS")
62     print("GANTT " + " ".join(gantt))
63     print(f"OK: AVG_WAIT {avg_wait:.2f} AVG_TAT {avg_tat:.2f}")
64
65 def sjf(processes):
66     time = 0
67     gantt = []
68     wait_times = []
69     tat_times = []
70     remaining = processes.copy()
71
72     while remaining:
73         available = [p for p in remaining if p.arrival <= time]
74         if not available:
75             next_arrival = min(remaining, key=lambda p: p.arrival)
76             gantt.append(f"IDLE@{time}-{next_arrival.arrival}")
77             time = next_arrival.arrival
78             available = [p for p in remaining if p.arrival <= time]
```

```
79     available.sort(key=lambda p: (p.burst, p.arrival, p.pid))
80     p = available[0]
81     start = time
82     end = start + p.burst
83     gantt.append(f"{p.pid}@{start}-{end}")
84     wait_times.append(start - p.arrival)
85     tat_times.append(end - p.arrival)
86     time = end
87     remaining.remove(p)
88
89     avg_wait = round(sum(wait_times)/len(wait_times), 2)
90     avg_tat = round(sum(tat_times)/len(tat_times), 2)
91
92     print("ALG SJF")
93     print("GANTT " + " ".join(gantt))
94     print(f"OK: AVG_WAIT {avg_wait:.2f} AVG_TAT {avg_tat:.2f}")
95
96 def main():
97     processes = parse_input()
98     fcfs(processes)
99     sjf(processes)
100
101 if __name__ == "__main__":
102     main()
```

Sample Input

```
$ printf 'pid,arrival,burst
P1,0,5
P2,2,2
P3,4,1
' | python cpu_scheduling.py
```

Sample Output

```
ALG FCFS
GANTT P1@0-5 P2@5-7 P3@7-8
OK: AVG_WAIT 2.00 AVG_TAT 4.67
ALG SJF
GANTT P1@0-5 P3@5-6 P2@6-8
OK: AVG_WAIT 1.67 AVG_TAT 4.33
```

Experiment 12: CPU Scheduling Simulator II (Round Robin)

Problem Statement

- Implement a CLI tool named `schedsim2`.
- Input is provided via `stdin` as CSV with header: `pid,arrival,burst`.
- Argument: `--q <quantum>` (integer 1..1000).
- Simulate preemptive Round Robin scheduling:
 - Newly arrived processes are enqueued at the end at their arrival time.
 - When a time slice ends and the running process is not finished, enqueue it at the end.
 - If CPU becomes idle, time jumps to the next process arrival.
- Output exactly:

```
ALG RR
GANTT <pid1>@<t0>-<t1> <pid2>@<t1>-<t2> ...
OK: AVG_WAIT <w> AVG_TAT <t>
```

Include `IDLE` segments if CPU is idle.

- On error, output one `ERROR:` line and exit non-zero.

Solution Implementation

```
1
2 import sys
3 import csv
4 import argparse
5
```

```
6 def error(code, msg):
7     print(f"ERROR: {code}: {msg}")
8     sys.exit(1)
9
10 # Parse arguments
11 parser = argparse.ArgumentParser()
12 parser.add_argument("--q", type=int, required=True)
13 args = parser.parse_args()
14
15 quantum = args.q
16 if not (1 <= quantum <= 1000):
17     error("E_RANGE", "quantum must be in 1..1000")
18
19 # Read CSV from stdin
20 try:
21     reader = csv.DictReader(sys.stdin)
22     processes = []
23     pids_set = set()
24     for row in reader:
25         pid = row["pid"]
26         if pid in pids_set:
27             error("E_DUPPID", f"duplicate pid {pid}")
28         pids_set.add(pid)
29         try:
30             arrival = int(row["arrival"])
31             burst = int(row["burst"])
32             if burst <= 0:
33                 error("E_RANGE", "burst must be positive")
34         except:
35             error("E_INPUT", "invalid arrival or burst")
36         processes.append({
37             "pid": pid,
38             "arrival": arrival,
39             "burst": burst,
40             "remaining": burst,
41             "completion": 0
42         })
43 except Exception:
```

```
44     error("E_INPUT", "malformed CSV")
45
46 # Sort by arrival time, then PID lexicographically
47 processes.sort(key=lambda x: (x["arrival"], x["pid"]))
48
49 time = 0
50 ready_queue = []
51 gantt = []
52
53 # Keep track of processes left
54 remaining_processes = processes.copy()
55
56 while remaining_processes or ready_queue:
57     # Add new arrivals at current time
58     arrivals = [p for p in remaining_processes if p["arrival"] <= time]
59     arrivals.sort(key=lambda x: x["pid"]) # tie-breaker
60     for p in arrivals:
61         ready_queue.append(p)
62         remaining_processes.remove(p)
63
64     if not ready_queue:
65         if remaining_processes:
66             next_arrival = min(remaining_processes, key=lambda x: x[
67                 "arrival"])
68             gantt.append(f"IDLE@{time}-{next_arrival['arrival']}")
69             time = next_arrival["arrival"]
70             continue
71         else:
72             break
73
74     # Pick first process in ready queue
75     current = ready_queue.pop(0)
76     run_time = min(current["remaining"], quantum)
77     gantt.append(f"{current['pid']}@{time}-{time + run_time}")
78     time += run_time
79     current["remaining"] -= run_time
80
81     # Check if current finished
```

```
81     if current["remaining"] == 0:
82         current["completion"] = time
83     else:
84         # Requeue unfinished process
85         ready_queue.append(current)
86
87 # Compute metrics
88 total_wt = 0
89 total_tat = 0
90 for p in processes:
91     tat = p["completion"] - p["arrival"]
92     wt = tat - p["burst"]
93     total_wt += wt
94     total_tat += tat
95
96 avg_wt = total_wt / len(processes)
97 avg_tat = total_tat / len(processes)
98
99 # Output
100 print("ALG RR")
101 print("GANTT", " ".join(gantt))
102 print(f"OK: AVG_WAIT {avg_wt:.2f} AVG_TAT {avg_tat:.2f}")
```

Sample Input

```
printf 'pid,arrival,burst
P1,0,5
P2,1,3
P3,2,8
P4,3,6' | python3 cpu_scheduling_round_robin.py --q 2
```

Sample Output

```
ALG RR
GANTT P1@0-2 P1@2-4 P2@4-6 P3@6-8 P1@8-9
P4@9-11 P2@11-12 P3@12-14 P4@14-16 P3@16-18 P4@18-20 P3@20-22
OK: AVG_WAIT 8.75 AVG_TAT 14.25
```

Experiment 13: Priority Scheduling Simulator (Non-preemptive with Aging)

Problem Statement

- Implement a CLI tool named `schedprio`.
- Input is provided via `stdin` as CSV with header: `pid,arrival,burst,priority`.
- Smaller priority means higher priority.
- Scheduling rules:
 - Non-preemptive: once a process starts, it runs to completion.
 - Aging: for every unit of waiting in the ready queue, effective priority decreases by 1 (minimum 0).
 - At dispatch, select ready process with lowest effective priority; tie-break by arrival then pid.
- Output exactly:

```
ALG PRIO_AGING
GANTT <pid1>@<t0>-<t1> <pid2>@<t1>-<t2> ...
OK: AVG_WAIT <w> AVG_TAT <t>
```

- On error, output one `ERROR:` line and exit non-zero.

Solution Implementation

```
1 #!/usr/bin/env python3
2 import sys
3
4 def error(code, msg):
```

```
5     print(f"ERROR: {code}: {msg}")
6     sys.exit(1)
7
8 def read_processes_from_stdin():
9     lines = [line.strip() for line in sys.stdin if line.strip()]
10    if not lines:
11        error("E_INPUT", "no input provided")
12    # Skip header if present
13    if lines[0].lower().startswith("pid"):
14        lines = lines[1:]
15    processes = []
16    seen_pid = set()
17    for i, line in enumerate(lines):
18        parts = line.split(",")
19        if len(parts) != 4:
20            error("E_INPUT", f"invalid input format in line {i+2}")
21        pid = parts[0].strip()
22        if pid in seen_pid:
23            error("E_DUPPID", f"duplicate pid '{pid}'")
24        seen_pid.add(pid)
25        try:
26            arrival = int(parts[1])
27            burst = int(parts[2])
28            priority = int(parts[3])
29        except:
30            error("E_INPUT", f"non-integer field in line {i+2}")
31        if arrival < 0:
32            error("E_RANGE", f"arrival must be >=0 in line {i+2}")
33        if burst <= 0:
34            error("E_RANGE", f"burst must be >0 in line {i+2}")
35        if not (0 <= priority <= 99):
36            error("E_RANGE", f"priority must be in 0..99 in line {i+2}")
37        processes.append({
38            "pid": pid,
39            "arrival": arrival,
40            "burst": burst,
41            "priority": priority
42        }
```

```
42         }
43     return processes
44
45 def schedule(processes):
46     time = 0
47     ready = []
48     todo = processes[:]
49     completed = {}
50     gantt = []
51
52     while todo or ready:
53         for p in todo[:]:
54             if p["arrival"] <= time:
55                 ready.append(p)
56                 todo.remove(p)
57
58             if not ready:
59                 next_time = min(p["arrival"] for p in todo)
60                 gantt.append(f"IDLE@{time}-{next_time}")
61                 time = next_time
62                 continue
63
64             # Aging at dispatch
65             for p in ready:
66                 waiting = time - p["arrival"]
67                 p["eff_prio"] = max(0, p["priority"] - waiting)
68
69             ready.sort(key=lambda x: (x["eff_prio"], x["arrival"], x["pid"]))
70             cur = ready.pop(0)
71
72             start = time
73             time += cur["burst"]
74             gantt.append(f"{cur['pid']}@{start}-{time}")
75             completed[cur["pid"]] = time
76
77             # Metrics
78             total_wait = 0
```

```
79     total_tat = 0
80     n = len(processes)
81
82     for p in processes:
83         ct = completed[p["pid"]]
84         tat = ct - p["arrival"]
85         wt = tat - p["burst"]
86         total_wait += wt
87         total_tat += tat
88
89     return gantt, total_wait/n, total_tat/n
90
91 def main():
92     processes = read_processes_from_stdin()
93     gantt, avg_wait, avg_tat = schedule(processes)
94     print("\nALG PRIO AGING")
95     print("GANTT", " ".join(gantt))
96     print(f"OK: AVG_WAIT {avg_wait:.2f} AVG_TAT {avg_tat:.2f}")
97
98 if __name__ == "__main__":
99     main()
```

Sample Input

```
printf 'pid,arrival,burst,priority
A,0,6,7
B,1,11,12
' | python3 priority_scheduling_cli.py
```

Sample Output

```
ALG PRIO_AGING
GANTT A@0-6 B@6-17
OK: AVG_WAIT 2.50 AVG_TAT 11.00
```

Experiment 14: Deadlock Avoidance using Banker's Algorithm

Problem Statement

- Implement a CLI tool named `banker`.
- Input is provided via `stdin` in the exact format:
 - First line: $P \ R$ (number of processes and resource types).
 - Next P lines: Allocation matrix (R integers per line).
 - Next P lines: Max matrix (R integers per line).
 - Last line: Available vector (R integers).
- Validate that $\text{Allocation}[i][j] \leq \text{Max}[i][j]$ for all entries.
- Apply Banker's safety algorithm.
- Output exactly:

OK: SAFE

OK: SEQ < p_0 > < p_1 > ... < $p_{(P-1)}$ >
- If the system is unsafe, output exactly:

OK: UNSAFE
- On error, output one `ERROR:` line and exit non-zero.

Solution Implementation

```
1 #!/usr/bin/env python3
2 import sys
```

```
3
4 def error(msg):
5     print(f"ERROR: {msg}")
6     sys.exit(1)
7
8 def read_input():
9     lines = [line.strip() for line in sys.stdin if line.strip()]
10    if len(lines) < 3:
11        error("Not enough input lines")
12    try:
13        P, R = map(int, lines[0].split())
14    except:
15        error("First line must be 'P R' with integers")
16
17    expected_lines = 1 + P*2 + 1
18    if len(lines) != expected_lines:
19        error(f"Expected {expected_lines} lines for {P} processes and {R} resources, got {len(lines)}")
20
21    # Allocation matrix
22    allocation = []
23    for i in range(P):
24        try:
25            row = list(map(int, lines[1 + i].split()))
26        except:
27            error(f"Invalid integers in Allocation row {i}")
28        if len(row) != R:
29            error(f"Allocation row {i} must have {R} entries")
30        allocation.append(row)
31
32    # Max matrix
33    max_need = []
34    for i in range(P):
35        try:
36            row = list(map(int, lines[1 + P + i].split()))
37        except:
38            error(f"Invalid integers in Max row {i}")
39        if len(row) != R:
```

```
40             error(f"Max row {i} must have {R} entries")
41         max_need.append(row)
42
43     # Available vector
44     try:
45         available = list(map(int, lines[-1].split()))
46     except:
47         error("Invalid integers in Available vector")
48     if len(available) != R:
49         error(f"Available vector must have {R} entries")
50
51     # Basic validation
52     for i in range(P):
53         for j in range(R):
54             if allocation[i][j] > max_need[i][j]:
55                 error(f"Allocation cannot exceed Max for process {i},"
56                       "resource {j}")
57             if allocation[i][j] < 0 or max_need[i][j] < 0 or available[
58                 j] < 0:
59                 error("Resource values must be non-negative")
60
61     return allocation, max_need, available
62
63
64 def bankers_algorithm(allocation, max_need, available):
65     n = len(allocation)
66     m = len(available)
67     need = [[max_need[i][j] - allocation[i][j] for j in range(m)] for i
68             in range(n)]
69     finish = [False]*n
70     safe_sequence = []
71     work = available[:]
72
73     while len(safe_sequence) < n:
74         found = False
75         for i in range(n):
76             if not finish[i] and all(need[i][j] <= work[j] for j in
77                                     range(m)):
78                 for j in range(m):
```

```
74             work[j] += allocation[i][j]
75             finish[i] = True
76             safe_sequence.append(i)
77             found = True
78         if not found:
79             print("OK: UNSAFE")
80             return
81
82     print("OK: SAFE")
83     print("OK: SEQ", *safe_sequence)
84
85 def main():
86     allocation, max_need, available = read_input()
87     bankers_algorithm(allocation, max_need, available)
88
89 if __name__ == "__main__":
90     main()
```

Sample Input

```
printf '5 3
0 1 0
2 0 0
3 0 2
2 1 1
0 0 2
7 5 3
3 2 2
9 0 2
2 2 2
4 3 3
3 3 2
' | python3 deadlock_banker_algo.py
```

Sample Output

```
OK: SAFE
```

```
OK: SEQ 1 3 4 0 2
```

Experiment 15: Deadlock Detection via Wait-For Graph Cycle

Problem Statement

- Implement a CLI tool named `wfgcheck`.
- Input is provided via `stdin` in the following format:
 - First line: $P \ E$ (number of processes and edges).
 - Next E lines: $u \ v$ meaning process u waits for v (directed edge $u \rightarrow v$).
 - Processes are numbered $0..P-1$.
- Output exactly one of:
 - No deadlock:
OK: DEADLOCK NO
 - Deadlock detected:
OK: DEADLOCK YES
OK: CYCLE <p₀> <p₁> ... <p_k> <p₀>
- If multiple cycles exist, output the cycle with:
 - Smallest starting node.
 - If tied, smallest lexicographic sequence.
- On error, output one `ERROR:` line and exit non-zero.

Solution Implementation

```
1 #!/usr/bin/env python3
2 import sys
```

```
3 from collections import deque
4
5 def error(msg):
6     print(f"ERROR: {msg}")
7     sys.exit(1)
8
9 def read_input():
10    lines = [line.strip() for line in sys.stdin if line.strip()]
11    if len(lines) < 1:
12        error("No input provided")
13
14    try:
15        P, E = map(int, lines[0].split())
16    except:
17        error("First line must be two integers: P E")
18
19    if len(lines) != 1 + E:
20        error(f"Expected {E} edge lines, got {len(lines)-1}")
21
22    edges = []
23    for i, line in enumerate(lines[1:]):
24        parts = line.split()
25        if len(parts) != 2:
26            error(f"Edge line {i+1} must have 2 integers: '{line}'")
27        try:
28            u, v = map(int, parts)
29        except:
30            error(f"Cannot parse edge line {i+1}: '{line}'")
31        if not (0 <= u < P) or not (0 <= v < P):
32            error(f"Edge ({u},{v}) out of range 0..{P-1}")
33        edges.append((u, v))
34
35    return P, edges
36
37 def build_graph(P, edges):
38    graph = [[] for _ in range(P)]
39    indegree = [0]*P
40    for u, v in edges:
```

```
41         graph[u].append(v)
42         indegree[v] += 1
43     return graph, indegree
44
45 def kahn_deadlock(P, graph, indegree):
46     q = deque([i for i in range(P) if indegree[i] == 0])
47     processed = 0
48     indegree_copy = indegree.copy()
49     while q:
50         node = q.popleft()
51         processed += 1
52         for neigh in graph[node]:
53             indegree_copy[neigh] -= 1
54             if indegree_copy[neigh] == 0:
55                 q.append(neigh)
56     if processed == P:
57         return None
58
59     remaining = [i for i in range(P) if indegree_copy[i] > 0]
60     visited = [False]*P
61     rec_stack = [False]*P
62     canonical_cycle = None
63
64     def dfs(node, path):
65         nonlocal canonical_cycle
66         visited[node] = True
67         rec_stack[node] = True
68         path.append(node)
69         for neigh in sorted(graph[node]):
70             if neigh not in remaining:
71                 continue
72             if rec_stack[neigh]:
73                 idx = path.index(neigh)
74                 cycle = path[idx:] + [neigh]
75                 if canonical_cycle is None:
76                     canonical_cycle = cycle
77                 else:
78                     old_start = min(canonical_cycle[:-1])
```

```
79             new_start = min(cycle[:-1])
80             if new_start < old_start:
81                 canonical_cycle = cycle
82             elif new_start == old_start and cycle <
83                 canonical_cycle:
84                 canonical_cycle = cycle
85             elif not visited[neigh]:
86                 dfs(neigh, path.copy())
87             rec_stack[node] = False
88
89     for u in sorted(remaining):
90         if not visited[u]:
91             dfs(u, [])
92
93     return canonical_cycle
94
95 def main():
96     P, edges = read_input()
97     graph, indegree = build_graph(P, edges)
98     cycle = kahn_deadlock(P, graph, indegree)
99
100    if cycle:
101        print("OK: DEADLOCK YES")
102        print("OK: CYCLE", ' '.join(map(str, cycle)))
103    else:
104        print("OK: DEADLOCK NO")
105
106 if __name__ == "__main__":
107     main()
```

Sample Input 1

```
printf '4 4  
0 1  
1 2  
2 0  
2 3  
' | python3 wfgcheck.py
```

Sample Output 1

```
OK: DEADLOCK YES  
OK: CYCLE 0 1 2 0
```

Sample Input 2

```
printf '4 3  
0 1  
1 2  
2 3  
' | python3 wfgcheck.py
```

Sample Output 2

```
OK: DEADLOCK NO
```

Experiment 16: Contiguous Memory Allocation Simulator (First/Best/Worst Fit)

Problem Statement

- Implement a CLI tool named memfit.
- Input is provided via `stdin` in the exact format:
 - Line 1: `B` — number of memory blocks.
 - Line 2: `B` integers — block sizes.
 - Line 3: `P` — number of processes.
 - Line 4: `P` integers — process sizes.
- Simulate the following algorithms independently, each starting from the original block list:
 - FIRST_FIT
 - BEST_FIT
 - WORST_FIT
- For each algorithm, output exactly:

```
ALG <name>
PROC <i> SIZE <s> -> BLOCK <j>
PROC <i> SIZE <s> -> FAIL
OK: ALLOCATED <k>/<P>
```
- On error, output one `ERROR:` line and exit non-zero.

Solution Implementation

```
2 import sys
3
4 def error(msg):
5     print(f"ERROR: {msg}")
6     sys.exit(1)
7
8 def read_input():
9     try:
10         lines = [line.strip() for line in sys.stdin if line.strip()]
11         if len(lines) < 4:
12             error("Not enough input lines")
13
14         B = int(lines[0])
15         if B <= 0:
16             error("B must be positive")
17
18         block_sizes = list(map(int, lines[1].split()))
19         if len(block_sizes) != B or any(b <= 0 for b in block_sizes):
20             error("Block sizes invalid or non-positive")
21
22         P = int(lines[2])
23         if P <= 0:
24             error("P must be positive")
25
26         process_sizes = list(map(int, lines[3].split()))
27         if len(process_sizes) != P or any(p <= 0 for p in process_sizes):
28             error("Process sizes invalid or non-positive")
29
30         return block_sizes, process_sizes
31     except Exception as e:
32         error(f"E_INPUT: {e}")
33
34 def first_fit(blocks, processes):
35     allocation = []
36     blocks_copy = blocks.copy()
37     allocated_count = 0
38     for i, p_size in enumerate(processes):
```

```
39         placed = False
40         for j, b_size in enumerate(blocks_copy):
41             if b_size >= p_size:
42                 allocation.append(f"PROC {i} SIZE {p_size} -> BLOCK {j}")
43                 blocks_copy[j] -= p_size
44                 allocated_count += 1
45                 placed = True
46                 break
47             if not placed:
48                 allocation.append(f"PROC {i} SIZE {p_size} -> FAIL")
49         return allocation, allocated_count
50
51 def best_fit(blocks, processes):
52     allocation = []
53     blocks_copy = blocks.copy()
54     allocated_count = 0
55     for i, p_size in enumerate(processes):
56         best_index = -1
57         best_size = None
58         for j, b_size in enumerate(blocks_copy):
59             if b_size >= p_size and (best_size is None or b_size <
60             best_size):
61                 best_index = j
62                 best_size = b_size
63             if best_index != -1:
64                 allocation.append(f"PROC {i} SIZE {p_size} -> BLOCK {
65                     best_index}")
66                 blocks_copy[best_index] -= p_size
67                 allocated_count += 1
68             else:
69                 allocation.append(f"PROC {i} SIZE {p_size} -> FAIL")
70     return allocation, allocated_count
71
72 def worst_fit(blocks, processes):
73     allocation = []
74     blocks_copy = blocks.copy()
75     allocated_count = 0
```

```
74     for i, p_size in enumerate(processes):
75         worst_index = -1
76         worst_size = None
77         for j, b_size in enumerate(blocks_copy):
78             if b_size >= p_size and (worst_size is None or b_size >
79             worst_size):
80                 worst_index = j
81                 worst_size = b_size
82         if worst_index != -1:
83             allocation.append(f"PROC {i} SIZE {p_size} -> BLOCK {
84             worst_index}")
85             blocks_copy[worst_index] -= p_size
86             allocated_count += 1
87     else:
88         allocation.append(f"PROC {i} SIZE {p_size} -> FAIL")
89     return allocation, allocated_count
90
91
92
93
94
95
96
97
98
99
100
101
102
103
104
105
106
107
```

```
def worst_fit(blocks, processes):
    allocation, allocated_count = first_fit(blocks, processes)
    print(f"\nALG WORST-FIT")
    for line in allocation:
        print(line)
    print(f"OK: ALLOCATED {allocated_count}/{len(processes)}")
```

```
def main():
    blocks, processes = read_input()
    run_all_algorithms(blocks, processes)
```

```
if __name__ == "__main__":
    main()
```

Sample Input

```
printf '5
100 500 200 300 600
4
212 417 112 426
' | python3 mem_allocation.py
```

Sample Output

```
ALG FIRST_FIT
PROC 0 SIZE 212 -> BLOCK 1
PROC 1 SIZE 417 -> BLOCK 4
PROC 2 SIZE 112 -> BLOCK 1
PROC 3 SIZE 426 -> FAIL
OK: ALLOCATED 3/4

ALG BEST_FIT
PROC 0 SIZE 212 -> BLOCK 3
PROC 1 SIZE 417 -> BLOCK 1
PROC 2 SIZE 112 -> BLOCK 2
PROC 3 SIZE 426 -> BLOCK 4
OK: ALLOCATED 4/4

ALG WORST_FIT
PROC 0 SIZE 212 -> BLOCK 4
PROC 1 SIZE 417 -> BLOCK 1
PROC 2 SIZE 112 -> BLOCK 4
PROC 3 SIZE 426 -> FAIL
OK: ALLOCATED 3/4
```

Experiment 17: Paging Address Translation with Optional TLB

Problem Statement

- Implement a CLI tool named pagetrans.
- Command-line arguments:
 - `--pagesize <S>` where S is a power of two in the range 256..65536.
 - `--tlb <K>` where K is the TLB size (0..64).
- Input via `stdin` in the following format:
 - Line 1: N — number of page table entries.
 - Next N lines: vpn pfn valid (valid is 0 or 1).
 - Next line: Q — number of queries.
 - Next Q lines: vaddr (unsigned decimal).
- For each query, output exactly one line:
 - Valid mapping:

OK: VA <vaddr> → PA <paddr> (TLB HIT|TLB MISS)
when K>0; omit the TLB part when K=0.
 - Page fault:

OK: VA <vaddr> → PAGEFAULT
- After all queries, if K>0, output exactly:

OK: TLB_HITS <h> TLB_MISSES <m>
- On error, output one `ERROR:` line and exit non-zero.

Solution Implementation

```
1 #!/usr/bin/env python3
2 import sys
3 import argparse
4
5 # --- Parse command-line arguments ---
6 parser = argparse.ArgumentParser()
7 parser.add_argument("--pagesize", type=int, required=True, help="Page
     size (must be power of 2)")
8 parser.add_argument("--tlb", type=int, required=True, help="TLB size
     (0..64)")
9 args = parser.parse_args()
10
11 pagesize = args.pagesize
12 tlb_size = args.tlb
13
14 # --- Validate arguments ---
15 if pagesize < 256 or pagesize > 65536 or (pagesize & (pagesize - 1)) != 0:
16     print(f"ERROR: E_RANGE: pagesize must be power of 2 between 256 and
           65536")
17     sys.exit(1)
18
19 if tlb_size < 0 or tlb_size > 64:
20     print(f"ERROR: E_RANGE: TLB size must be 0..64")
21     sys.exit(1)
22
23 # --- Read input from stdin ---
24 lines = [line.strip() for line in sys.stdin if line.strip()]
25 try:
26     # Number of page table entries
27     n = int(lines[0])
28     page_table = {}
29
30     # Read page table
31     for i in range(1, n+1):
32         parts = lines[i].split()
```

```
33     if len(parts) != 3:
34         raise ValueError(f"Invalid page table entry on line {i+1}")
35     vpn, pfn, valid = map(int, parts)
36     if valid not in (0, 1):
37         print(f"ERROR: E_INPUT: valid must be 0 or 1")
38         sys.exit(1)
39     page_table[vpn] = (pfn, valid)
40
41 # Number of queries
42 q_index = n + 1
43 Q = int(lines[q_index])
44 queries = [int(v) for v in lines[q_index+1 : q_index+1+Q]]
45 for v in queries:
46     if v < 0:
47         print(f"ERROR: E_RANGE: virtual address cannot be negative")
48         sys.exit(1)
49
50 except Exception as e:
51     print(f"ERROR: E_INPUT: {e}")
52     sys.exit(1)
53
54 # --- Initialize TLB if needed ---
55 TLB = [None] * tlb_size if tlb_size > 0 else None
56 tlb_hits = 0
57 tlb_misses = 0
58
59 # --- Process queries ---
60 for va in queries:
61     vpn = va // pagesize
62     offset = va % pagesize
63
64     # Page table check
65     if vpn not in page_table or page_table[vpn][1] == 0:
66         print(f"OK: VA {va} -> PAGEFAULT")
67         continue
68
69     pfn = page_table[vpn][0]
```

```
70
71     # TLB logic
72     if tlb_size > 0:
73         index = vpn % tlb_size
74         entry = TLB[index]
75         if entry is not None and entry[0] == vpn:
76             # TLB hit
77             tlb_hits += 1
78             print(f"OK: VA {va} -> PA {pfn*pagesize + offset} (TLB HIT)
79             ")
80     else:
81         # TLB miss
82         tlb_misses += 1
83         TLB[index] = (vpn, pfn)
84         print(f"OK: VA {va} -> PA {pfn*pagesize + offset} (TLB MISS
85             ")
86     else:
87         # No TLB
88         print(f"OK: VA {va} -> PA {pfn*pagesize + offset}")
89
90 # --- Print TLB stats ---
91 if tlb_size > 0:
92     print(f"OK: TLB_HITS {tlb_hits} TLB_MISSES {tlb_misses}")
```

Sample Input

```
printf '4
0 5 1
1 9 1
2 3 0
3 7 1
5
0
256
512
768
1024
' | python3 pagetrans.py --pagesize 512 --tlb 4
```

Sample Output

```
OK: VA 0 -> PA 2560 (TLB MISS)
OK: VA 256 -> PA 2816 (TLB HIT)
OK: VA 512 -> PA 4608 (TLB MISS)
OK: VA 768 -> PA 4864 (TLB HIT)
OK: VA 1024 -> PAGEFAULT
OK: TLB_HITS 2 TLB_MISSES 2
```

Experiment 18: Page Replacement Simulator (FIFO, LRU, OPT)

Problem Statement

- Implement a CLI tool named pagerepl.
- Command-line argument:
 - `--frames <F>` where F is the number of frames (1..64).
- Input via `stdin` in the following format:
 - Line 1: L — length of the reference string.
 - Line 2: L integers — page numbers (≥ 0).
- Simulate the following algorithms independently:
 - FIFO
 - LRU
 - OPT (Belady's optimal)
- For each algorithm, output exactly:

```
ALG <name>
OK: FAULTS <k>
OK: FINAL <f0> <f1> ... <f(F-1)>
```

where empty frames are printed as -1.
- On error, output one `ERROR:` line and exit non-zero.

Solution Implementation

```
1 # 18
2 import sys
```

```
3 import argparse
4 from typing import List
5
6 def print_result(alg: str, faults: int, frames: List[int]) -> None:
7     print(f"ALG {alg}")
8     print(f"OK: FAULTS {faults}")
9     print("OK: FINAL " + " ".join(str(f) if f != -1 else "-1" for f in
10                                    frames))
11
12 def simulate_fifo(refs: List[int], F: int) -> tuple[int, List[int]]:
13     """FIFO with circular pointer implementation."""
14     frames = [-1] * F
15     ptr = 0 # Pointer to next frame to replace
16     faults = 0
17     occupied = 0 # How many frames are actually occupied
18
19     for page in refs:
20         if page in frames:
21             continue # Hit
22
23         faults += 1 # Miss
24
25         if occupied < F:
26             # Fill empty frame
27             frames[occupied] = page
28             occupied += 1
29         else:
30             # Replace at pointer
31             frames[ptr] = page
32             ptr = (ptr + 1) % F
33
34     return faults, frames
35
36 def simulate_lru(refs: List[int], F: int) -> tuple[int, List[int]]:
37     """LRU with timestamp tracking."""
38     frames = [-1] * F
39     last_used = {} # page -> last access time
40     time = 0
```

```
40     faults = 0
41     occupied = 0
42
43     for page in refs:
44         if page in frames:
45             # Update last used time
46             last_used[page] = time
47         else:
48             faults += 1
49
50         if occupied < F:
51             # Fill first empty frame
52             frames[occupied] = page
53             last_used[page] = time
54             occupied += 1
55         else:
56             # Find LRU page in frames
57             lru_page = frames[0]
58             lru_time = last_used.get(lru_page, float('inf'))
59
60             for i in range(1, F):
61                 current_page = frames[i]
62                 current_time = last_used.get(current_page, float(
63                     'inf'))
64                 if current_time < lru_time:
65                     lru_page = current_page
66                     lru_time = current_time
67
68             # Replace LRU page
69             idx = frames.index(lru_page)
70             del last_used[lru_page]
71             frames[idx] = page
72             last_used[page] = time
73
74     time += 1
75
76     return faults, frames
```

```
77 def simulate_opt(refs: List[int], F: int) -> tuple[int, List[int]]:
78     """Optimal page replacement."""
79     frames = [-1] * F
80     faults = 0
81     occupied = 0
82
83     for i, page in enumerate(refs):
84         if page in frames:
85             continue
86
87         faults += 1
88
89         if occupied < F:
90             # Fill empty frame
91             frames[occupied] = page
92             occupied += 1
93         else:
94             # Find page to evict
95             victim_idx = 0
96             farthest_next = -1
97
98             for idx, frame_page in enumerate(frames):
99                 # Find next use of this page
100                next_use = len(refs) + 1 # Default: never used again
101                for j in range(i + 1, len(refs)):
102                    if refs[j] == frame_page:
103                        next_use = j
104                        break
105
106                if next_use == len(refs) + 1:
107                    # Never used again - evict this one immediately
108                    victim_idx = idx
109                    break
110                elif next_use > farthest_next:
111                    farthest_next = next_use
112                    victim_idx = idx
113                elif next_use == farthest_next and idx < victim_idx:
114                    # Tie-break: smaller frame index
```

```
115                 victim_idx = idx
116
117             frames[victim_idx] = page
118
119     return faults, frames
120
121 def main():
122     parser = argparse.ArgumentParser()
123     parser.add_argument('--frames', type=int, required=True)
124     args = parser.parse_args()
125
126     if args.frames < 1 or args.frames > 64:
127         print("ERROR: E_RANGE: frames must be 1..64", file=sys.stderr)
128         return 1
129
130     try:
131         L_line = sys.stdin.readline()
132         if not L_line:
133             return 1
134         L = int(L_line.strip())
135
136         refs_line = sys.stdin.readline()
137         if not refs_line:
138             return 1
139         refs = list(map(int, refs_line.strip().split()))
140
141         if len(refs) != L:
142             print("ERROR: E_INPUT: length mismatch", file=sys.stderr)
143             return 1
144
145         if any(x < 0 for x in refs):
146             print("ERROR: E_RANGE: page numbers must be >= 0", file=sys
147 .stderr)
147             return 1
148     except Exception:
149         print("ERROR: E_INPUT: invalid input", file=sys.stderr)
150         return 1
151
```

```
152     # Run all algorithms
153     f_faults, f_frames = simulate_fifo(refs, args.frames)
154     print_result("FIFO", f_faults, f_frames)
155
156     l_faults, l_frames = simulate_lru(refs, args.frames)
157     print_result("LRU", l_faults, l_frames)
158
159     o_faults, o_frames = simulate_opt(refs, args.frames)
160     print_result("OPT", o_faults, o_frames)
161
162     return 0
163
164 if __name__ == "__main__":
165     sys.exit(main())
```

Sample Input

```
printf '12
1 2 3 4 1 2 5 1 2 3 4 5
' | python3 pagerepl.py --frames 3
```

Sample Output

```
ALG FIFO
OK: FAULTS 9
OK: FINAL 5 3 4
ALG LRU
OK: FAULTS 10
OK: FINAL 3 4 5
ALG OPT
OK: FAULTS 7
OK: FINAL 4 2 5
```

Experiment 19: File Allocation Strategy Simulator (Contiguous, Linked, Indexed)

Problem Statement

- Implement a CLI tool named `filealloc`.
- Input is provided via `stdin` in the following format:
 - Line 1: `N` — total number of disk blocks (1..10000).
 - Line 2: `F` — number of free blocks.
 - Line 3: `F` integers — free block IDs (0..N-1, unique).
 - Line 4: `M` — number of files.
 - Next `M` lines: `name size`, where `size` is blocks required (1..N).
- Simulate the following allocation strategies independently, each starting from the original free list:
 - CONTIGUOUS
 - LINKED
 - INDEXED
- Output per algorithm:

```
ALG <name>
FILE <name> -> <map>
FILE <name> -> FAIL
```
- Map formats:
 - CONTIGUOUS: `START LEN <size>`
 - LINKED: `CHAIN <b1>-><b2>->...`
 - INDEXED: `INDEX <i> DATA <b1>, <b2>, ...`
- On error, output one `ERROR: line` and exit non-zero.

Solution Implementation

1 | # 19

```
2 import sys
3 from typing import List, Tuple, Optional, Set
4
5 def error_exit(code: str, message: str) -> None:
6     """Print error message and exit with non-zero status."""
7     print(f"ERROR: {code}: {message}", file=sys.stderr)
8     sys.exit(1)
9
10 def read_input() -> Tuple[int, Set[int], List[Tuple[str, int]]]:
11     """Read and validate input from stdin."""
12     try:
13         # Read total blocks
14         line = sys.stdin.readline()
15         if not line:
16             error_exit("E_INPUT", "empty input")
17         N = int(line.strip())
18         if not (1 <= N <= 10000):
19             error_exit("E_RANGE", "total blocks must be 1..10000")
20
21         # Read number of free blocks
22         line = sys.stdin.readline()
23         if not line:
24             error_exit("E_INPUT", "missing free blocks count")
25         F = int(line.strip())
26
27         # Read free block IDs
28         line = sys.stdin.readline()
29         if not line:
30             error_exit("E_INPUT", "missing free block list")
31         free_blocks = list(map(int, line.strip().split()))
32
33         if len(free_blocks) != F:
34             error_exit("E_INPUT", f"free block count mismatch: expected
35                         {F}, got {len(free_blocks)}")
36
37         # Validate free blocks
38         seen = set()
39         for block in free_blocks:
```

```
39         if not (0 <= block < N):
40             error_exit("E_RANGE", f"block ID {block} out of range
41                         0..{N-1}")
42             if block in seen:
43                 error_exit("E_DUPBLOCK", "free block list must contain
44                             unique IDs")
45             seen.add(block)
46
47             # Read number of files
48             line = sys.stdin.readline()
49             if not line:
50                 error_exit("E_INPUT", "missing file count")
51             M = int(line.strip())
52
53             # Read file requests
54             files = []
55             for _ in range(M):
56                 line = sys.stdin.readline()
57                 if not line:
58                     error_exit("E_INPUT", "incomplete file list")
59                 parts = line.strip().split()
60                 if len(parts) != 2:
61                     error_exit("E_INPUT", "invalid file line format")
62
63                 name = parts[0]
64                 try:
65                     size = int(parts[1])
66                 except ValueError:
67                     error_exit("E_INPUT", f"invalid file size for {name}")
68
69                 if not (1 <= size <= N):
70                     error_exit("E_RANGE", f"file size for {name} must be
71                         1..{N}")
72
73                 files.append((name, size))
74
75             return N, set(free_blocks), files
```

```
74     except ValueError:
75         error_exit("E_INPUT", "invalid number format")
76
77 def simulate_contiguous(N: int, free_blocks: Set[int], files: List[
78     Tuple[str, int]]) -> List[Tuple[str, str]]:
79     """Simulate contiguous allocation strategy."""
80     results = []
81     # Create a sorted list of free blocks for easier contiguous search
82     sorted_free = sorted(free_blocks)
83
84     for name, size in files:
85         allocated = False
86
87         # Find smallest starting block with enough consecutive free
88         # blocks
89         for i in range(len(sorted_free) - size + 1):
90             start = sorted_free[i]
91             # Check if blocks start..start+size-1 are all free
92             consecutive = True
93             for j in range(1, size):
94                 if sorted_free[i + j] != start + j:
95                     consecutive = False
96                     break
97
98             if consecutive:
99                 # Allocation successful
100                results.append((name, f"START {start} LEN {size}"))
101
102                # Remove allocated blocks from free list
103                for j in range(size):
104                    sorted_free.remove(start + j)
105
106                allocated = True
107                break
108
109                if not allocated:
110                    results.append((name, "FAIL"))
```

```
110     return results
111
112 def simulate_linked(N: int, free_blocks: Set[int], files: List[Tuple[
113     str, int]]) -> List[Tuple[str, str]]:
114     """Simulate linked allocation strategy."""
115     results = []
116     # Use sorted list for deterministic smallest block selection
117     available_blocks = sorted(free_blocks)
118
119     for name, size in files:
120         if len(available_blocks) >= size:
121             # Allocate smallest 'size' blocks
122             allocated = available_blocks[:size]
123
124             # Create chain string
125             chain = "->".join(str(b) for b in allocated)
126             results.append((name, f"CHAIN {chain}"))
127
128             # Remove allocated blocks
129             available_blocks = available_blocks[size:]
130         else:
131             results.append((name, "FAIL"))
132
133     return results
134
135 def simulate_indexed(N: int, free_blocks: Set[int], files: List[Tuple[
136     str, int]]) -> List[Tuple[str, str]]:
137     """Simulate indexed allocation strategy."""
138     results = []
139     # Use sorted list for deterministic smallest block selection
140     available_blocks = sorted(free_blocks)
141
142     for name, size in files:
143         # Indexed allocation needs: 1 index block + size data blocks
144         total_needed = size + 1
145
146         if len(available_blocks) >= total_needed:
```

```
145     # Allocate blocks: first block is index, next 'size' blocks
146     # are data
147     index_block = available_blocks[0]
148     data_blocks = available_blocks[1:size+1]
149
150     # Create data blocks string
151     data_str = ",".join(str(b) for b in data_blocks)
152     results.append((name, f"INDEX {index_block} DATA {data_str}")
153                     ))
154
155     # Remove allocated blocks
156     available_blocks = available_blocks[total_needed:]
157
158 else:
159     results.append((name, "FAIL"))
160
161 return results
162
163 def main() -> None:
164     """Main function implementing filealloc tool."""
165     # Read and validate input
166     N, free_blocks, files = read_input()
167
168     # Simulate each strategy independently
169     print("ALG CONTIGUOUS")
170     contiguous_results = simulate_contiguous(N, free_blocks.copy(),
171                                              files)
172     for name, result in contiguous_results:
173         print(f"FILE {name} -> {result}")
174
175     print("\nALG LINKED")
176     linked_results = simulate_linked(N, free_blocks.copy(), files)
177     for name, result in linked_results:
178         print(f"FILE {name} -> {result}")
179
180     print("\nALG INDEXED")
181     indexed_results = simulate_indexed(N, free_blocks.copy(), files)
182     for name, result in indexed_results:
183         print(f"FILE {name} -> {result}")
```

```
180  
181 if __name__ == "__main__":  
182     main()
```

Sample Input

```
printf '20  
10  
1 2 3 4 5 7 8 10 11 12  
3  
A 3  
B 4  
C 6  
' | python3 filealloc.py
```

Sample Output

```
ALG CONTIGUOUS  
FILE A -> START 1 LEN 3  
FILE B -> FAIL  
FILE C -> FAIL  
  
ALG LINKED  
FILE A -> CHAIN 1->2->3  
FILE B -> CHAIN 4->5->7->8  
FILE C -> FAIL  
  
ALG INDEXED  
FILE A -> INDEX 1 DATA 2,3,4  
FILE B -> INDEX 5 DATA 7,8,10,11  
FILE C -> FAIL
```

Experiment 20: Disk Scheduling Simulator (FCFS, SSTF, SCAN, C-SCAN)

Problem Statement

- Implement a CLI tool named `disksched`.
- Command-line arguments:
 - `--max <C>` — maximum cylinder number ($C \geq 1$).
 - `--start <S>` — starting head position ($0 \leq S \leq C$).
 - `--dir left|right` — initial scan direction (used for SCAN and C-SCAN).
- Input via `stdin` in the following format:
 - Line 1: `L` — number of disk requests.
 - Line 2: `L` integers — requested cylinder numbers ($0..C$).
- Simulate the following algorithms independently:
 - FCFS
 - SSTF
 - SCAN
 - C-SCAN
- For each algorithm, output exactly:

```
ALG <name>
OK: ORDER <c1> <c2> ... <cL>
OK: MOVES <m>
```
- On error, output one `ERROR:` line and exit non-zero.

Solution Implementation

```
1
2 import sys
3 import argparse
4 from typing import List
5
6 def error_exit(code: str, message: str):
7     print(f"ERROR: {code}: {message}", file=sys.stderr)
8     sys.exit(1)
9
10 class DiskScheduler:
11     def __init__(self, max_cylinder: int, start: int, direction: str):
12         self.max_cylinder = max_cylinder
13         self.start = start
14         self.direction = direction # 'left' or 'right'
15
16     def fcfs(self, requests: List[int]) -> tuple[List[int], int]:
17         """First-Come, First-Served."""
18         order = requests.copy()
19         moves = self._calculate_moves(order)
20         return order, moves
21
22     def sstf(self, requests: List[int]) -> tuple[List[int], int]:
23         """Shortest Seek Time First."""
24         order = []
25         remaining = requests.copy()
26         current = self.start
27
28         while remaining:
29             # Find closest request
30             closest_idx = 0
31             closest_dist = abs(remaining[0] - current)
32
33             for i in range(1, len(remaining)):
34                 dist = abs(remaining[i] - current)
35                 if dist < closest_dist or (dist == closest_dist and
36                     remaining[i] < remaining[closest_idx]):
```

```
36             closest_dist = dist
37             closest_idx = i
38
39             order.append(remaining[closest_idx])
40             current = remaining[closest_idx]
41             remaining.pop(closest_idx)
42
43             moves = self._calculate_moves(order)
44             return order, moves
45
46     def scan(self, requests: List[int]) -> tuple[List[int], int]:
47         """SCAN (Elevator) algorithm."""
48         order = []
49         current = self.start
50
51         if self.direction == 'right':
52             # Moving right first
53             right_requests = [r for r in requests if r >= current]
54             right_requests.sort()
55             order.extend(right_requests)
56
57             # Then move to max cylinder if needed
58             if right_requests:
59                 order.append(self.max_cylinder)
60
61             # Then serve left side
62             left_requests = [r for r in requests if r < current]
63             left_requests.sort(reverse=True)
64             order.extend(left_requests)
65
66         else: # left
67             # Moving left first
68             left_requests = [r for r in requests if r <= current]
69             left_requests.sort(reverse=True)
70             order.extend(left_requests)
71
72             # Then move to 0 if needed
73             if left_requests:
```

```
74         order.append(0)
75
76         # Then serve right side
77         right_requests = [r for r in requests if r > current]
78         right_requests.sort()
79         order.extend(right_requests)
80
81         # Remove duplicates of end cylinders if they weren't in
82         # original requests
83         # (We added them to simulate going to the end)
84         final_order = []
85         for cyl in order:
86             if cyl in requests or cyl == 0 or cyl == self.max_cylinder:
87                 final_order.append(cyl)
88             if cyl in requests:
89                 # Remove one occurrence from requests list
90                 idx = requests.index(cyl)
91                 requests.pop(idx)
92
93         moves = self._calculate_moves(final_order)
94         return final_order, moves
95
96     def cscan(self, requests: List[int]) -> tuple[List[int], int]:
97         """C-SCAN (Circular SCAN) algorithm."""
98         order = []
99         current = self.start
100
101         if self.direction == 'right':
102             # Moving right first
103             right_requests = [r for r in requests if r >= current]
104             right_requests.sort()
105             order.extend(right_requests)
106
107             # Go to max cylinder
108             if right_requests:
109                 order.append(self.max_cylinder)
110
111             # Jump to 0 (counts as movement)
```

```
111         order.append(0)

112

113     # Then serve remaining left side
114     left_requests = [r for r in requests if r < current]
115     left_requests.sort()
116     order.extend(left_requests)

117

118     else: # left
119         # Moving left first
120         left_requests = [r for r in requests if r <= current]
121         left_requests.sort(reverse=True)
122         order.extend(left_requests)

123

124     # Go to 0
125     if left_requests:
126         order.append(0)

127

128     # Jump to max cylinder (counts as movement)
129     order.append(self.max_cylinder)

130

131     # Then serve remaining right side
132     right_requests = [r for r in requests if r > current]
133     right_requests.sort(reverse=True)
134     order.extend(right_requests)

135

136     # Remove end cylinders if they weren't in original requests
137     final_order = []
138     temp_requests = requests.copy()
139     for cyl in order:
140         if cyl in temp_requests or cyl == 0 or cyl == self.
141             max_cylinder:
142                 final_order.append(cyl)
143                 if cyl in temp_requests:
144                     idx = temp_requests.index(cyl)
145                     temp_requests.pop(idx)

146     moves = self._calculate_moves(final_order)
147     return final_order, moves
```

```
148
149     def _calculate_moves(self, order: List[int]) -> int:
150         """Calculate total head movement for service order."""
151         total = 0
152         current = self.start
153         for cyl in order:
154             total += abs(cyl - current)
155             current = cyl
156         return total
157
158 def main():
159     parser = argparse.ArgumentParser()
160     parser.add_argument('--max', type=int, required=True)
161     parser.add_argument('--start', type=int, required=True)
162     parser.add_argument('--dir', choices=['left', 'right'], required=True)
163
164     args = parser.parse_args()
165
166     # Validate
167     if args.max < 1:
168         error_exit('E_RANGE', 'max cylinder must be >= 1')
169     if not (0 <= args.start <= args.max):
170         error_exit('E_RANGE', f'start must be 0..{args.max}')
171
172     # Read input
173     try:
174         line = sys.stdin.readline()
175         if not line:
176             error_exit('E_INPUT', 'empty input')
177         L = int(line.strip())
178
179         line = sys.stdin.readline()
180         if not line:
181             error_exit('E_INPUT', 'missing requests')
182         requests = list(map(int, line.strip().split()))
183
184         if len(requests) != L:
```

```
185         error_exit('E_INPUT', f'expected {L} requests, got {len(
186             requests)}')
187
187     for req in requests:
188         if not (0 <= req <= args.max):
189             error_exit('E_RANGE', f'request {req} out of range 0..{
190                 args.max}')
191
191 except ValueError:
192     error_exit('E_INPUT', 'invalid number format')
193
194 # Create scheduler and run algorithms
195 scheduler = DiskScheduler(args.max, args.start, args.dir)
196
197 # FCFS
198 order, moves = scheduler.fcfs(requests.copy())
199 print("ALG FCFS")
200 print(f"OK: ORDER {' '.join(map(str, order))}")
201 print(f"OK: MOVES {moves}")
202
203 # SSTF
204 order, moves = scheduler.sstf(requests.copy())
205 print("\nALG SSTF")
206 print(f"OK: ORDER {' '.join(map(str, order))}")
207 print(f"OK: MOVES {moves}")
208
209 # SCAN
210 order, moves = scheduler.scan(requests.copy())
211 print("\nALG SCAN")
212 print(f"OK: ORDER {' '.join(map(str, order))}")
213 print(f"OK: MOVES {moves}")
214
215 # C-SCAN
216 order, moves = scheduler.cscan(requests.copy())
217 print("\nALG C-SCAN")
218 print(f"OK: ORDER {' '.join(map(str, order))}")
219 print(f"OK: MOVES {moves}")
220
```

```
221     return 0  
222  
223 if __name__ == "__main__":  
224     sys.exit(main())
```

Sample Input

```
printf '8  
98 183 37 122 14 124 65 67  
' | python3 disksched.py --max 199 --start 53 --dir right
```

Sample Output

```
ALG FCFS  
OK: ORDER 98 183 37 122 14 124 65 67  
OK: MOVES 640  
  
ALG SSTF  
OK: ORDER 65 67 37 14 98 122 124 183  
OK: MOVES 236  
  
ALG SCAN  
OK: ORDER 65 67 98 122 124 183 199 37 14  
OK: MOVES 331  
  
ALG C-SCAN  
OK: ORDER 65 67 98 122 124 183 199 0 14 37  
OK: MOVES 382
```